Predicting potential loan defaults

What makes a loan risky?

I want a to buy a new house!

Loan Application

Credit
★★★★

Income
★★★

Term
★★★★★

Personal Info
★★★

Credit history explained

Did I pay previous loans on time?

Example: excellent, good, or fair

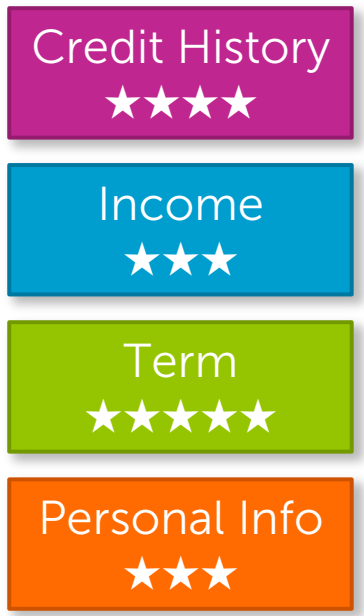

Income

What's my income?

Example:
$80K per year

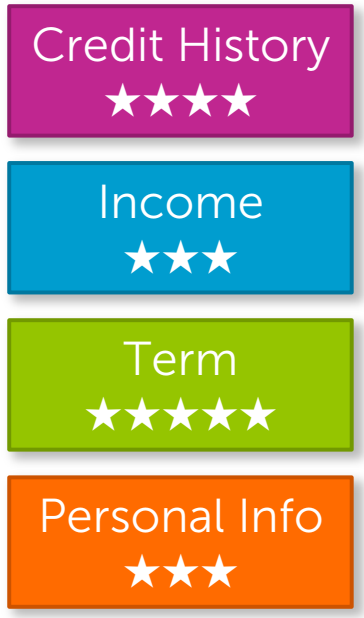

Loan terms

How soon do I need to pay the loan?

Example: 3 years, 5 years,...

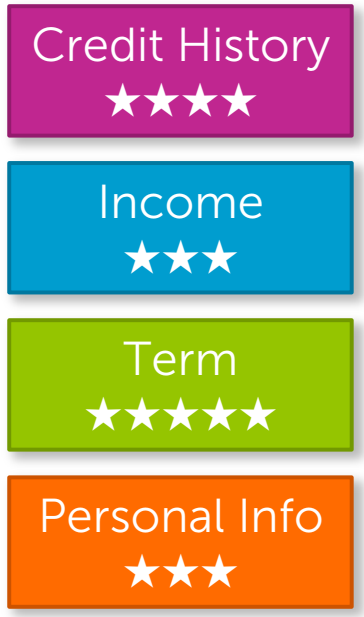

Personal information

Age, reason for the loan, marital status,...

Example: Home loan for a married couple

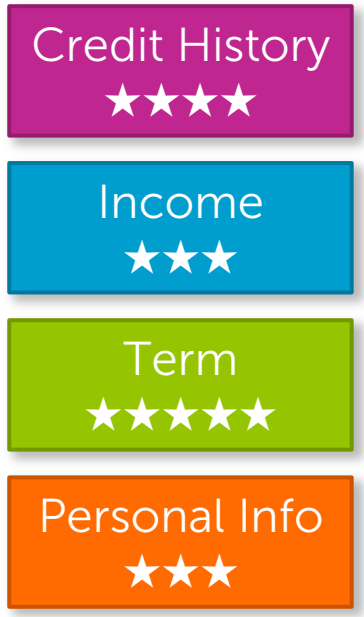

Intelligent application

Classifier review

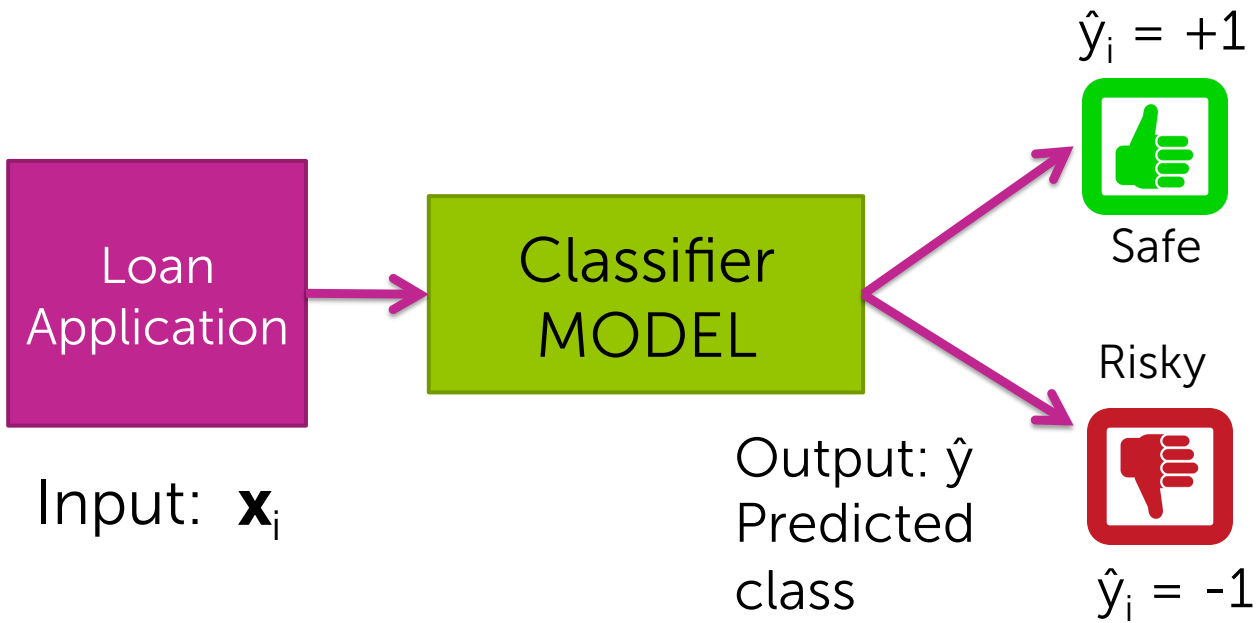

This module … decision trees

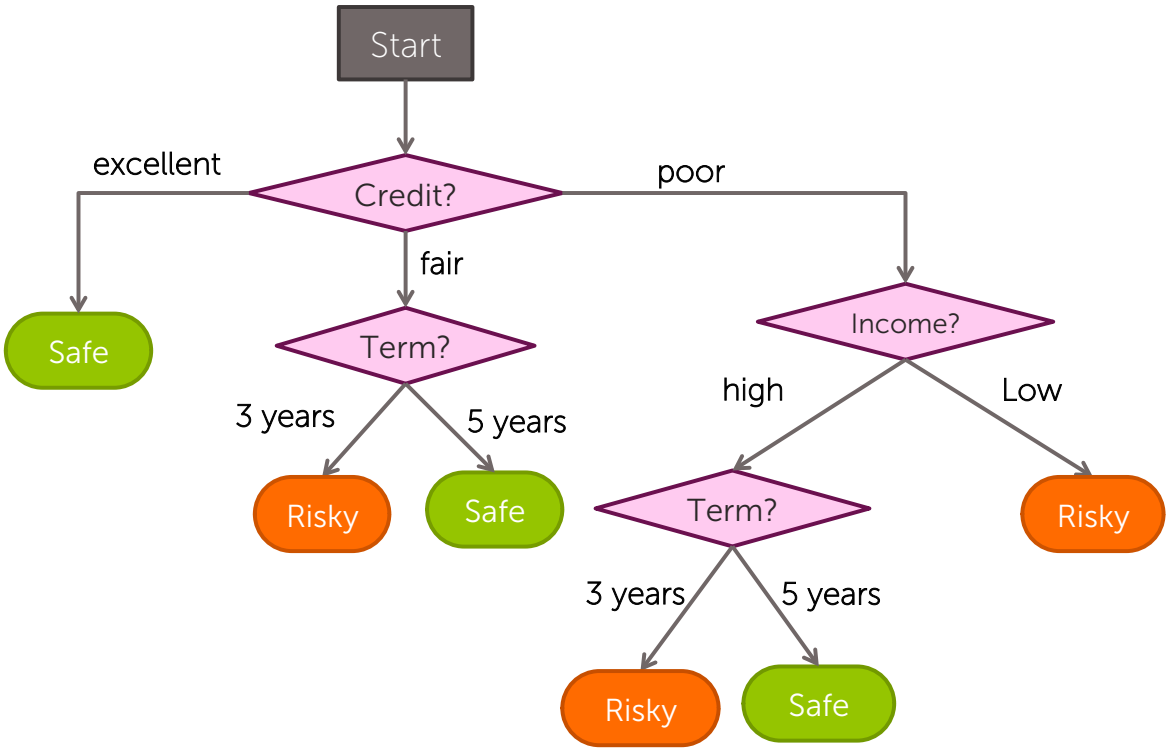

Decision trees: *Intuition*

What does a decision tree represent?

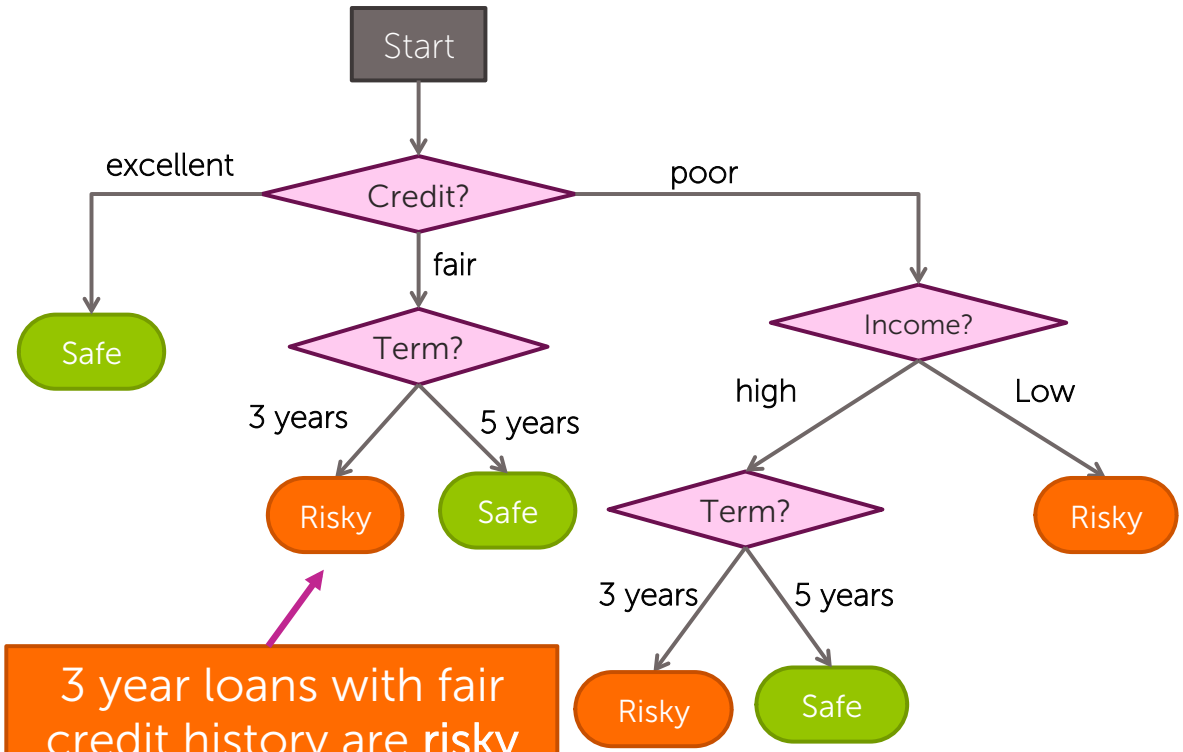

What does a decision tree represent?

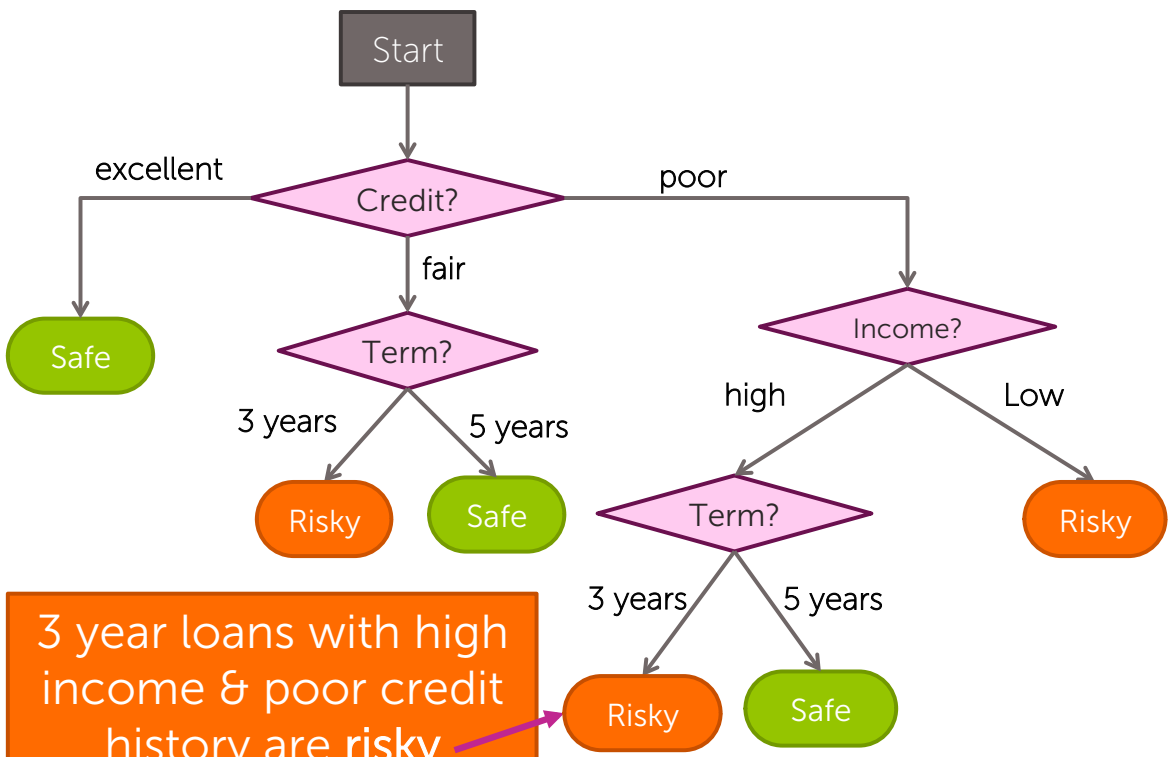

Scoring a loan application

$\mathbf{x}_i$ = (Credit = poor, Income = high, Term = 5 years)

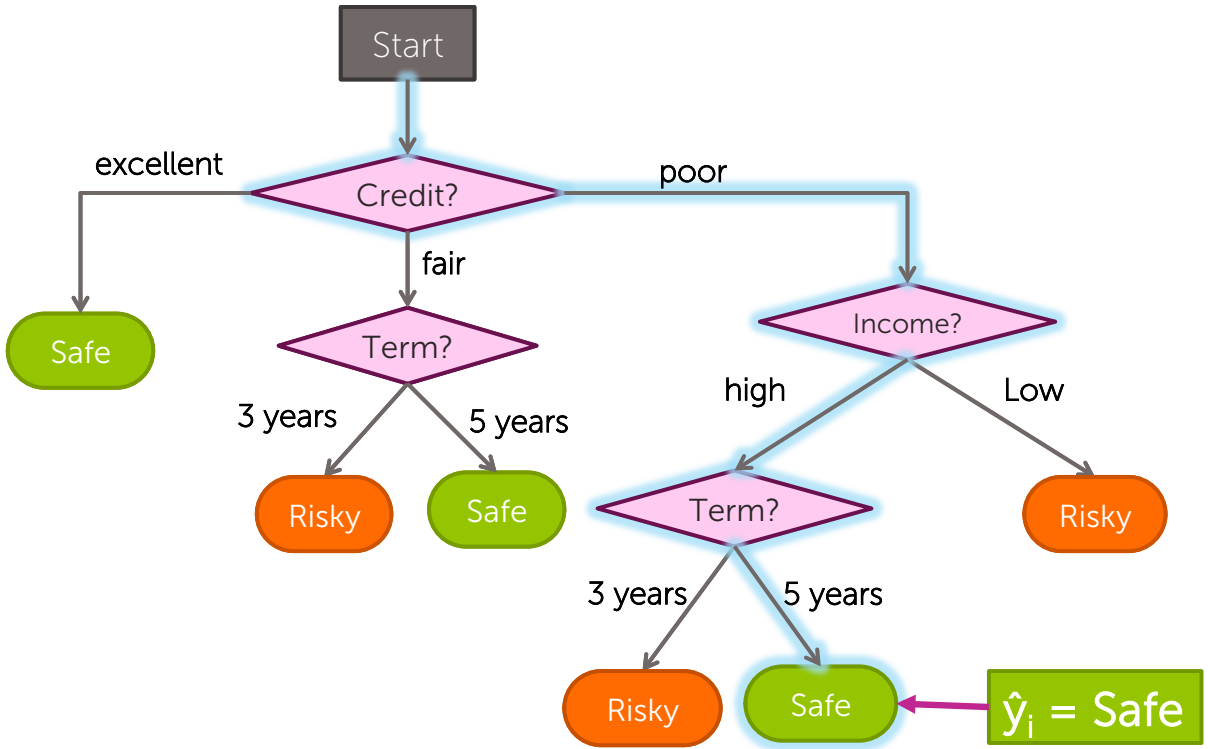

Decision tree model

Loan Application

Input: $\mathbf{x}_i$

$T(\mathbf{x}_i)$ = Traverse decision tree

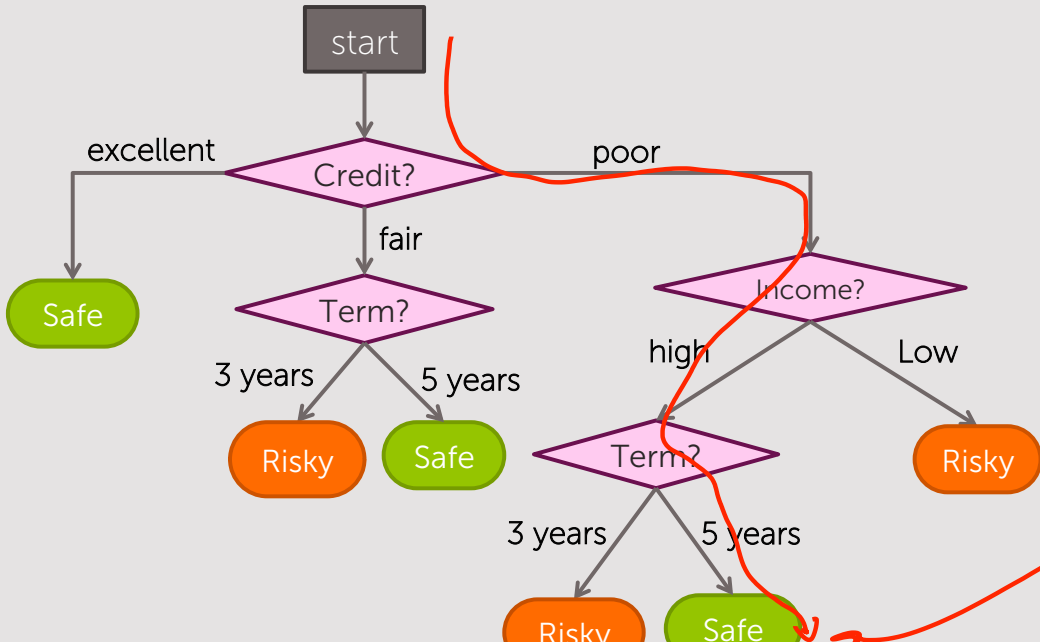

$\hat{y}_i$ = Safe

Decision tree learning task

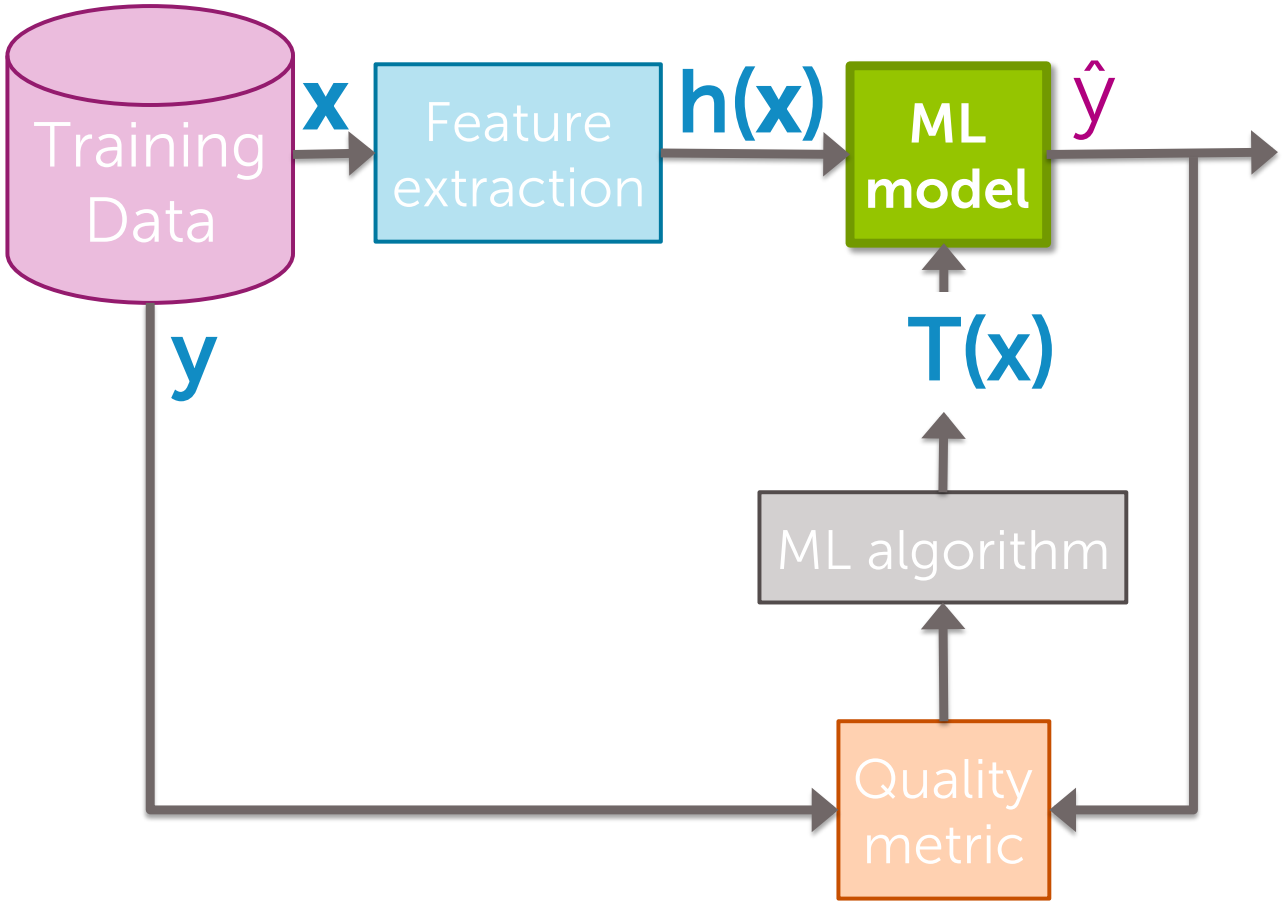
Training Data
x
Feature extraction
h(x)
ML model
ŷ
y
T(x)
ML algorithm
Quality metric

Learn decision tree from data?

$h_1(x)$ Credit	$h_2(x)$ Term	$h_3(x)$ Income	Loan Status y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

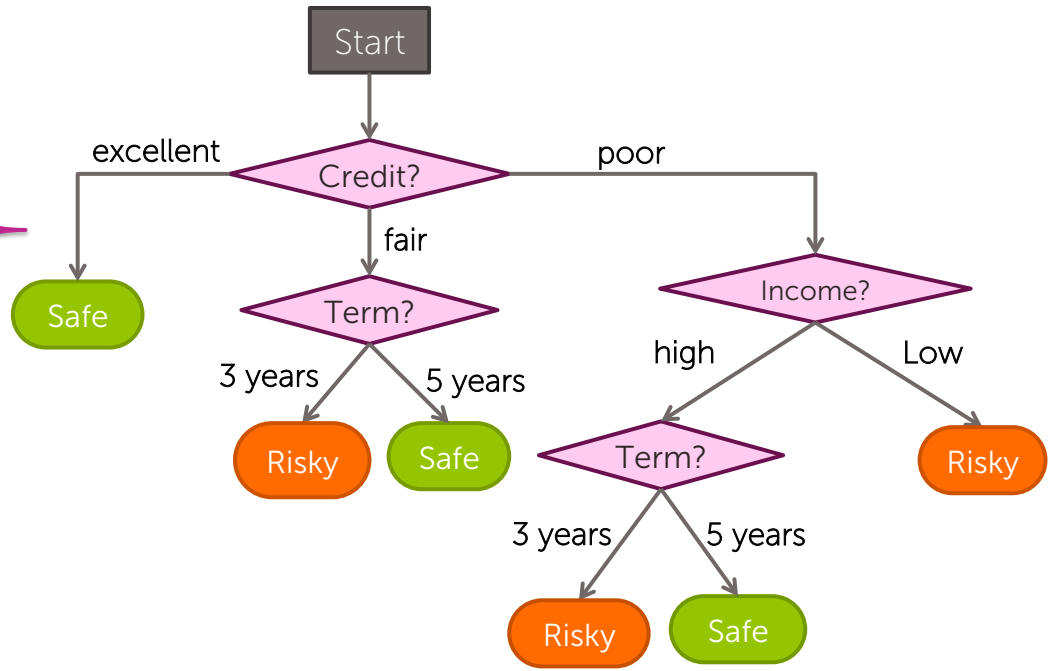

Decision tree learning problem

Training data: N observations ($\mathbf{x}_i$,y_i)

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

Optimize
quality metric
on training data

T(X)

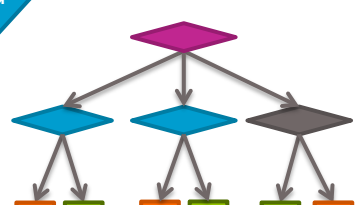

Quality metric: Classification error

- Error measures fraction of mistakes

$$\text{Error} = \frac{\#\text{ incorrect predictions}}{\#\text{ examples}}$$

  - Best possible value : 0.0
  - Worst possible value: 1.0

Find the tree with lowest classification error

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

Minimize
classification error
on training data

T(X)

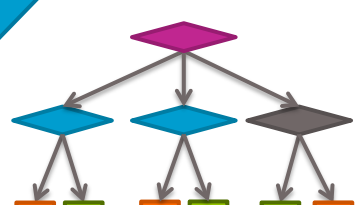

How do we find the best tree?

Exponentially large number of possible trees makes decision tree learning hard!
(NP-hard problem)

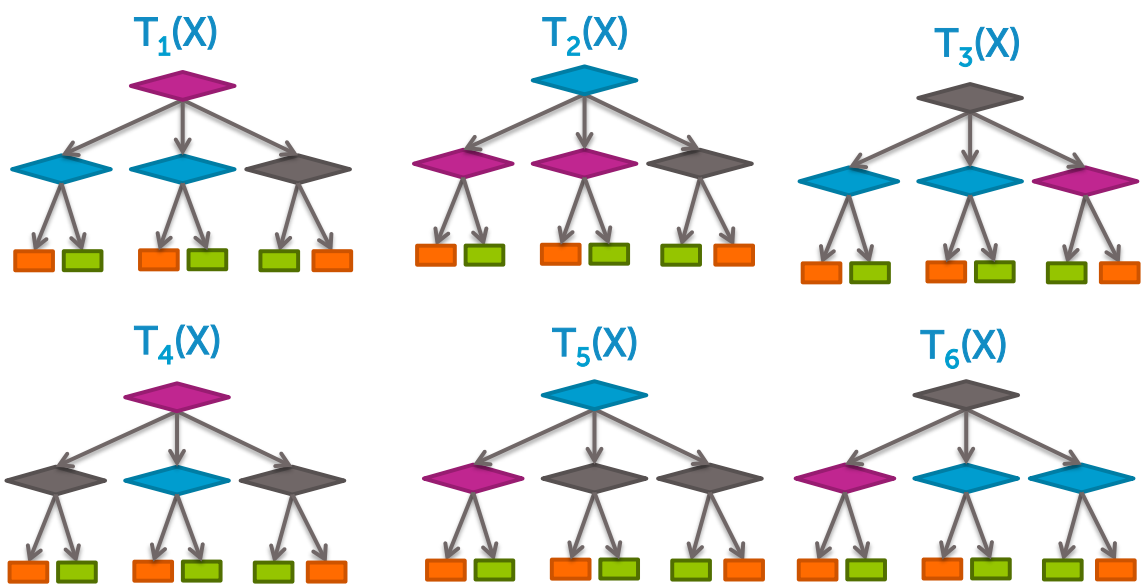

Simple (greedy) algorithm finds “good” tree

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

Approximately minimize **classification error** on training data

T(X)

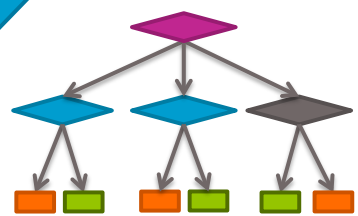

Greedy decision tree learning:
Algorithm outline

Step 1: Start with an empty tree

Histogram of y values

All points in the dataset

Step 2: Split on a feature

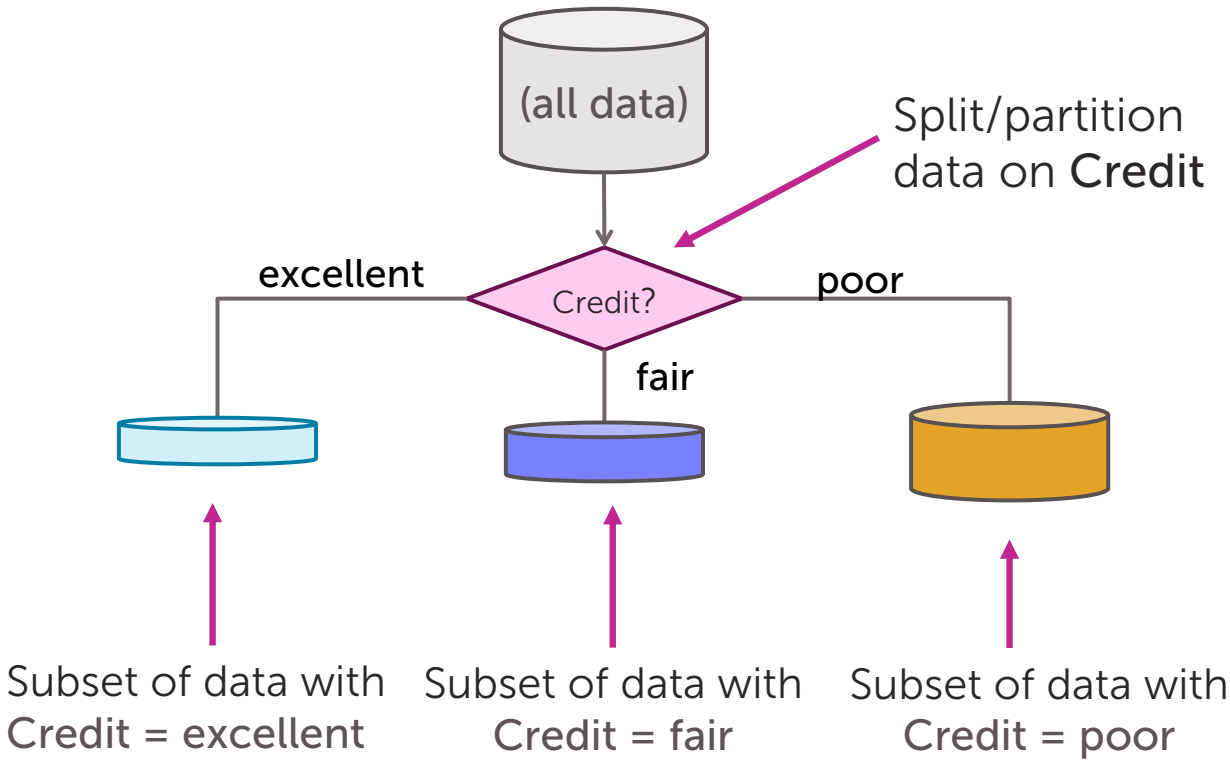

Feature split explained

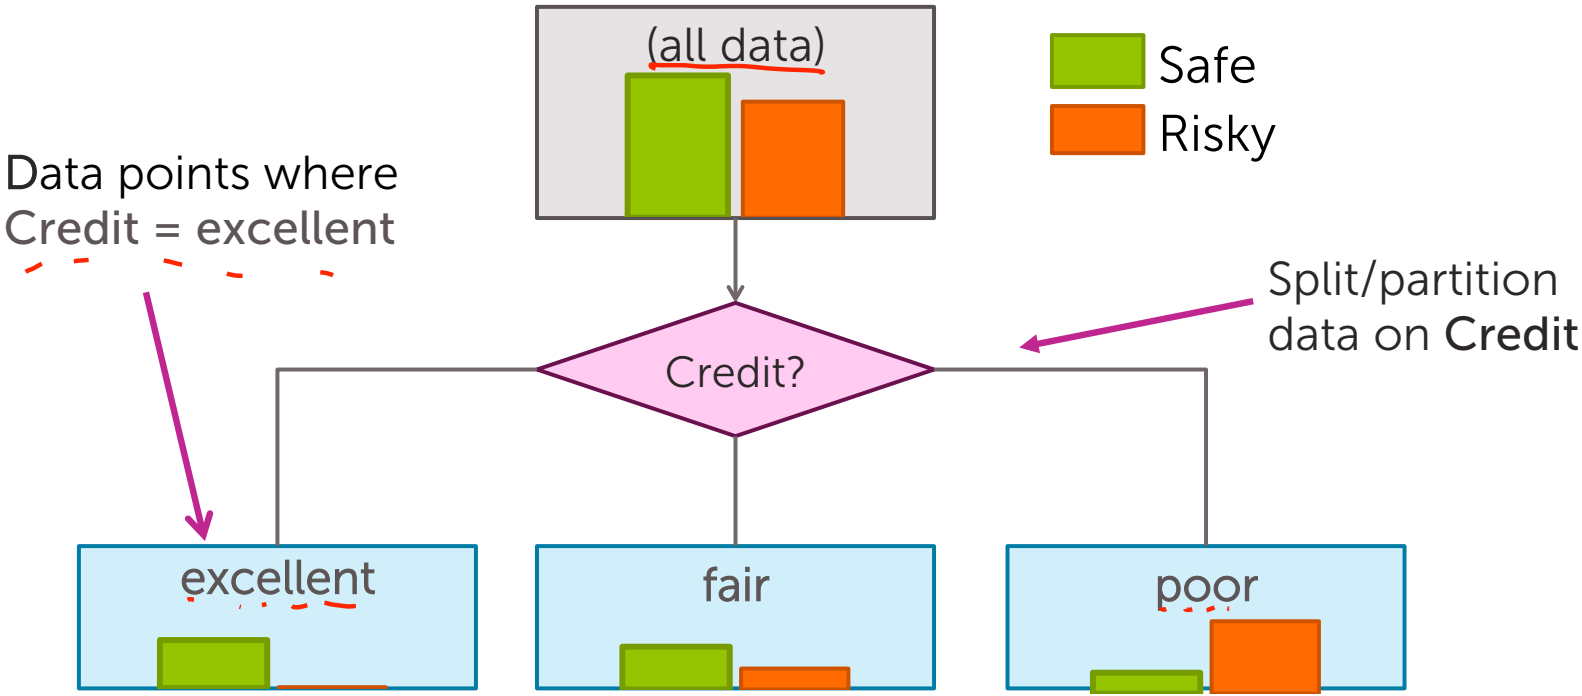

Step 3: Making predictions

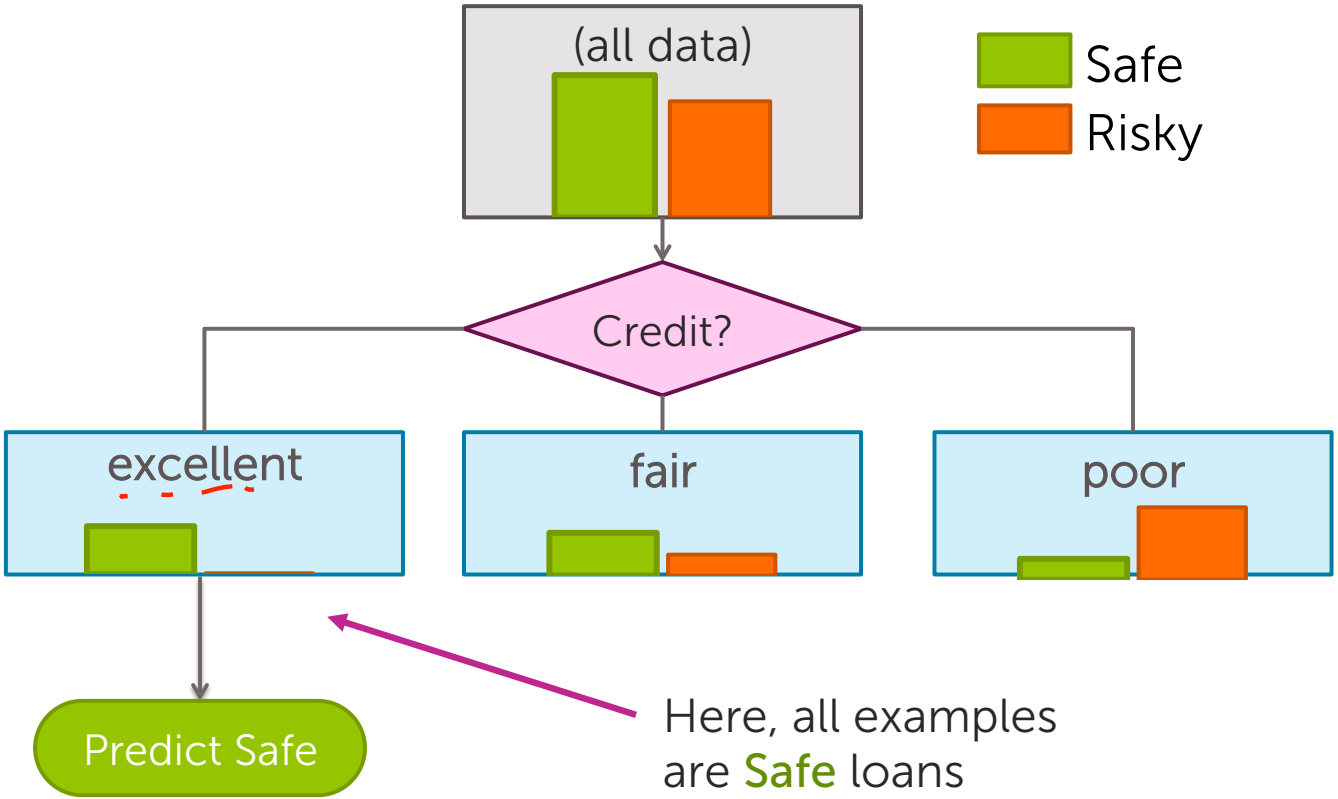

Step 4: Recursion

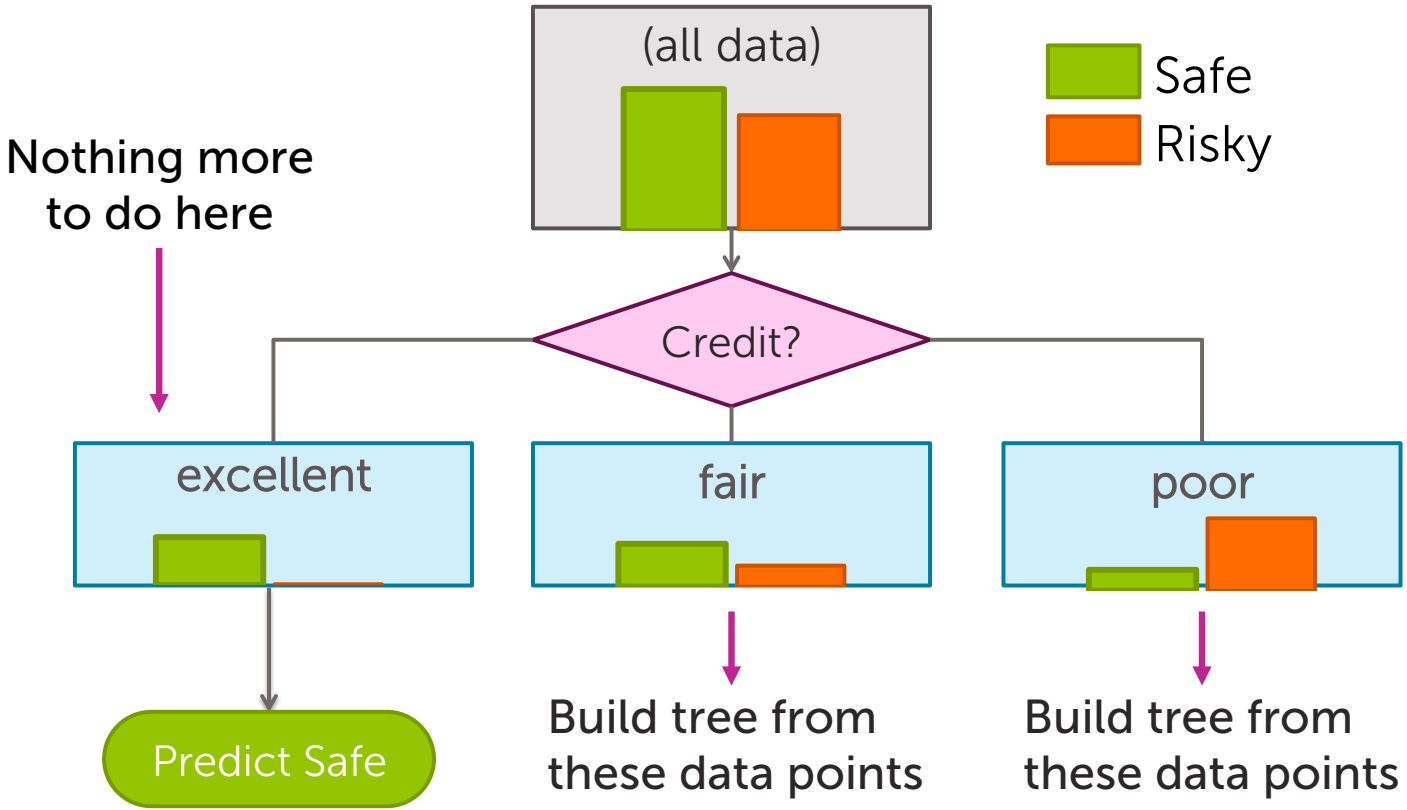

Greedy decision tree learning

- Step 1: Start with an empty tree
- Step 2: Select a feature to split data
- For each split of the tree:
  - Step 3: If nothing more to, make predictions
  - Step 4: Otherwise, go to Step 2 & continue (recurse) on this split

Problem 1: Feature split selection

Problem 2: Stopping condition

Recursion

Feature split learning

=

Decision stump learning

Start with the data

Assume N = 40, 3 features

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

Start with all the data

Loan status: Safe Risky

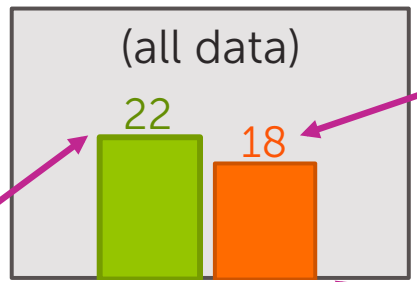

Number of Risky loans

Number of Safe loans

N = 40 examples

Compact visual notation: Root node

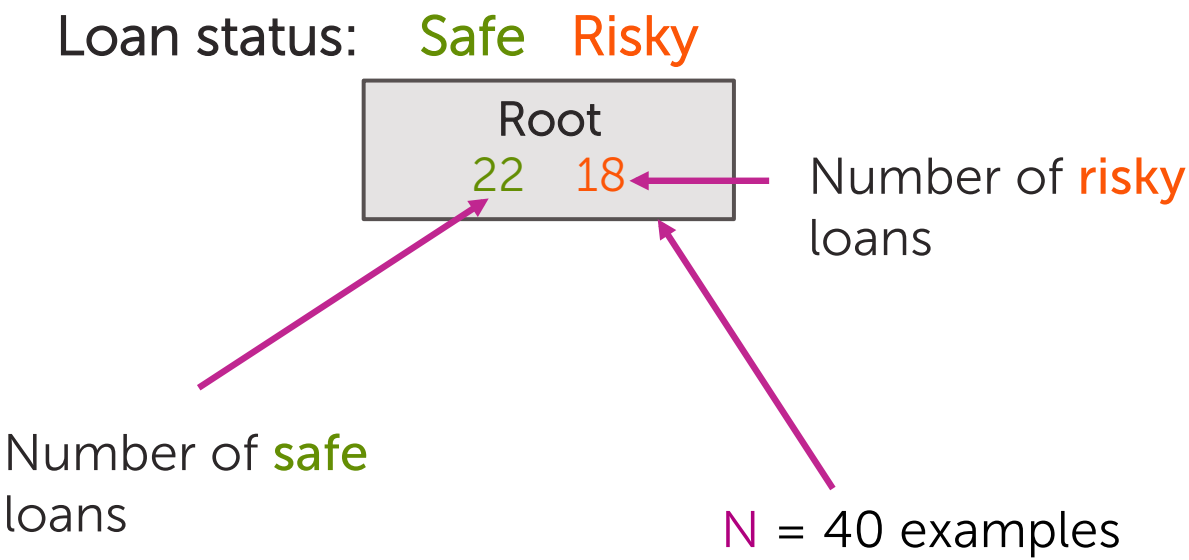

Decision stump: Single level tree

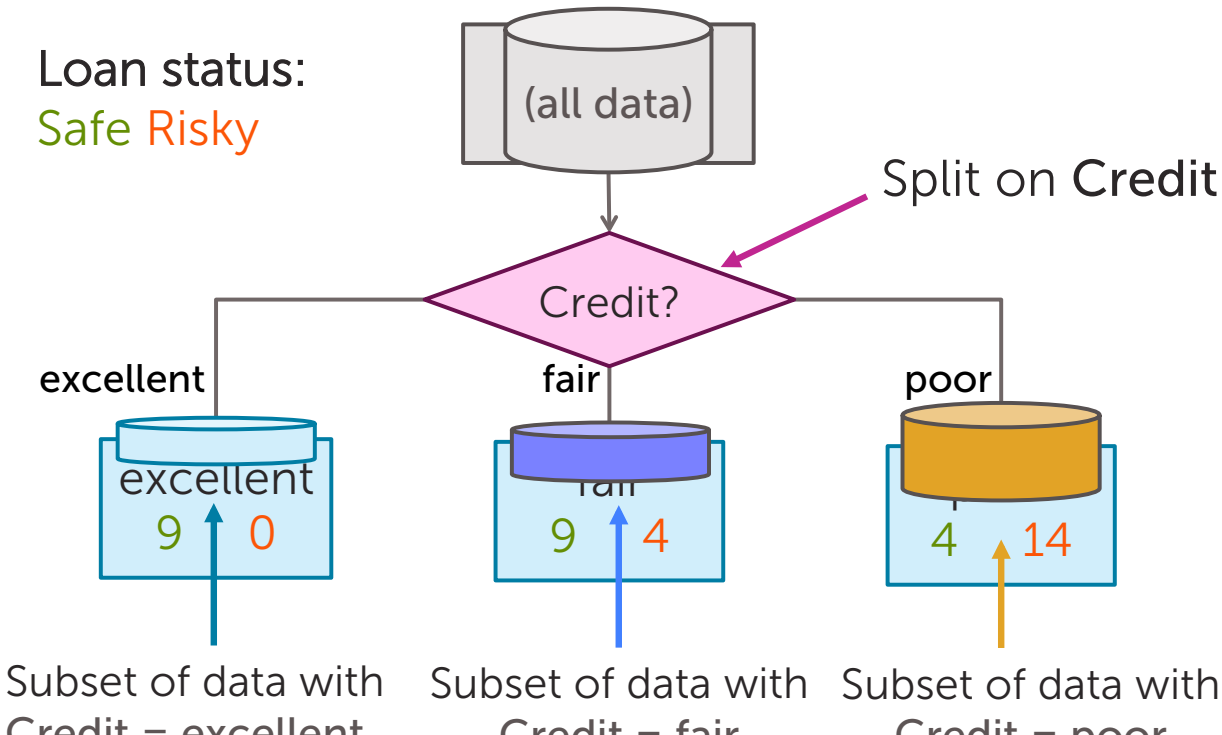

Visual Notation: Intermediate nodes

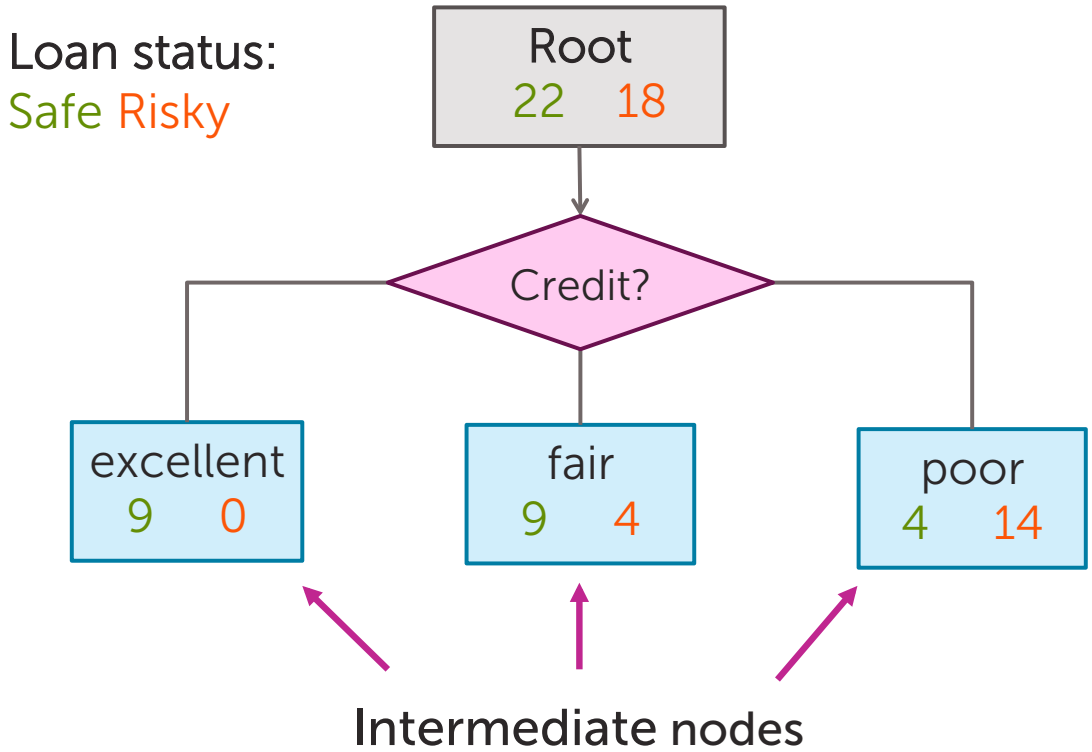

Making predictions with a decision stump

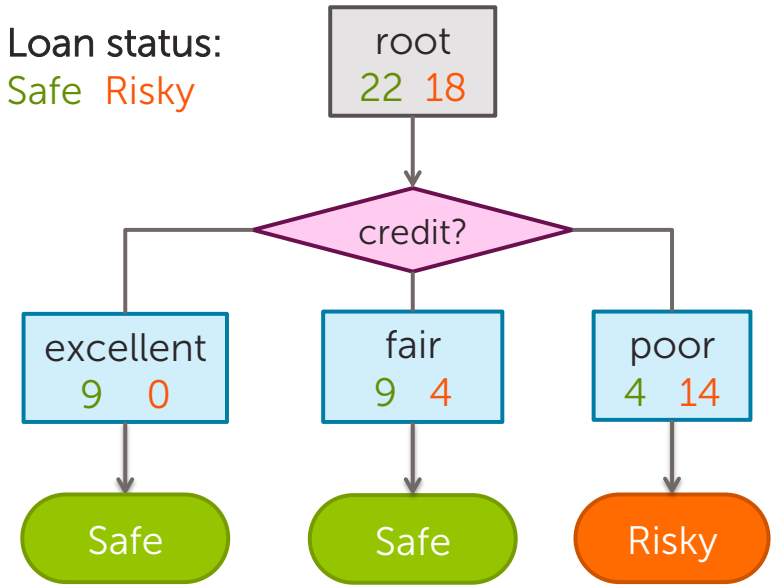

For each intermediate node, set $\hat{y}$ = **majority value**

Selecting best feature to split on

How do we learn a decision stump?

Loan status:
Safe Risky

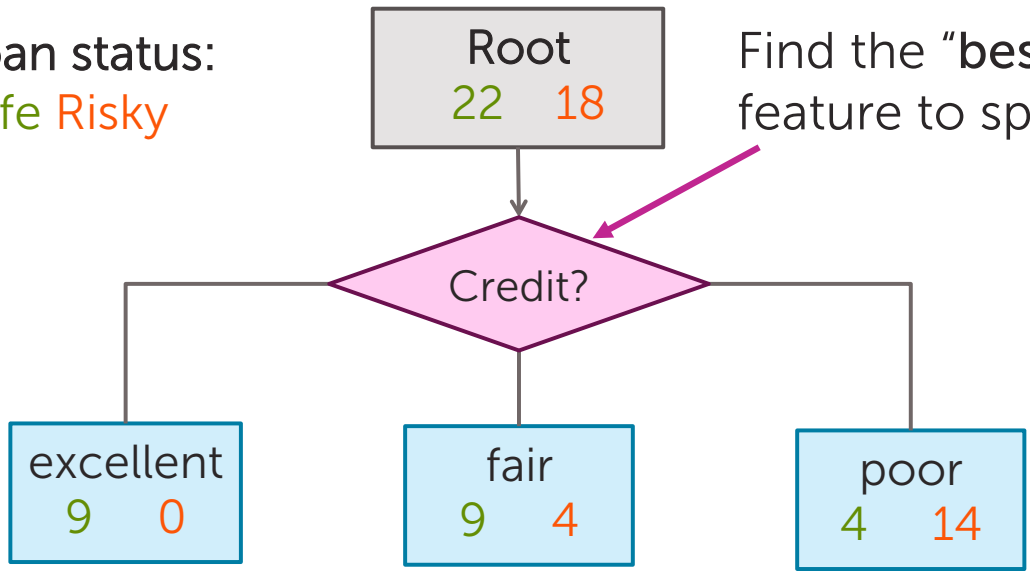

Find the "**best**" feature to split on!

How do we select the best feature?

Better?

Choice 1: Split on Credit

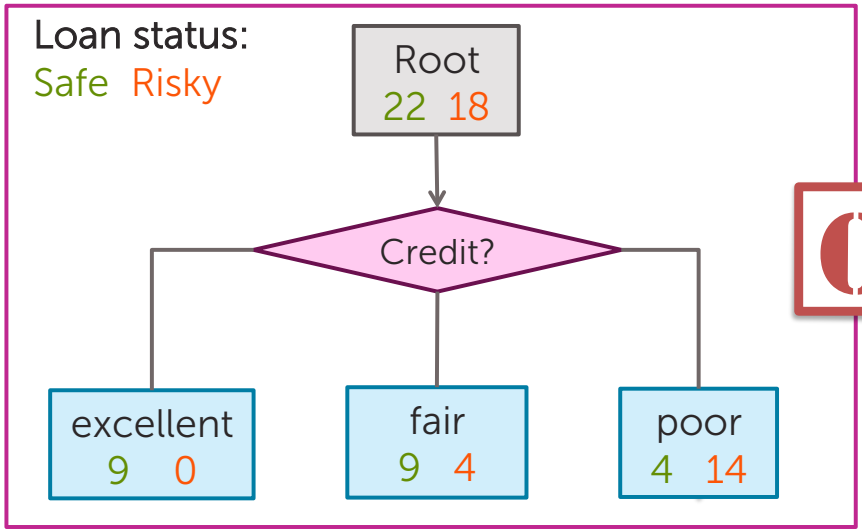

OR

Choice 2: Split on Term

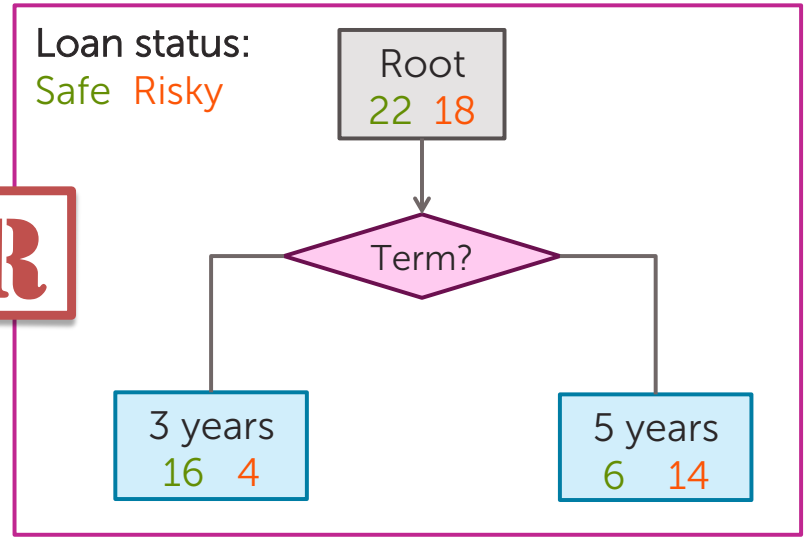

How do we measure effectiveness of a split?

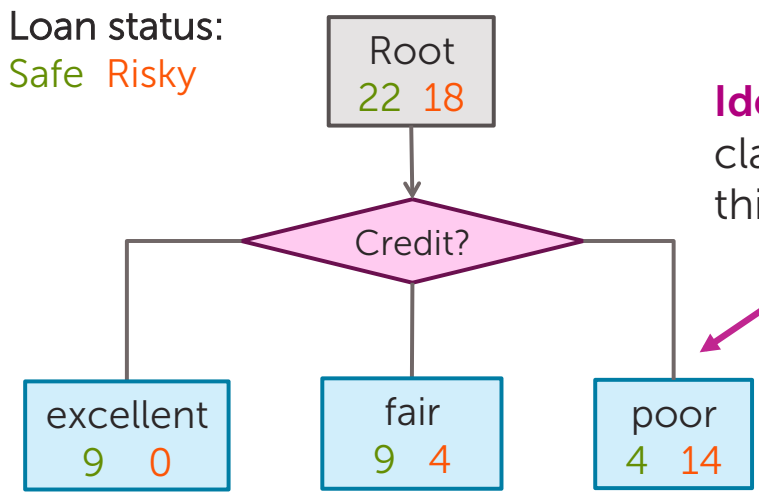

Idea: Calculate classification error of this decision stump

$$\text{Error} = \frac{\#\text{ mistakes}}{\#\text{ data points}}$$

Calculating classification error

- **Step 1:** ŷ = class of majority of data in node
- **Step 2:** Calculate classification error of predicting ŷ for this data

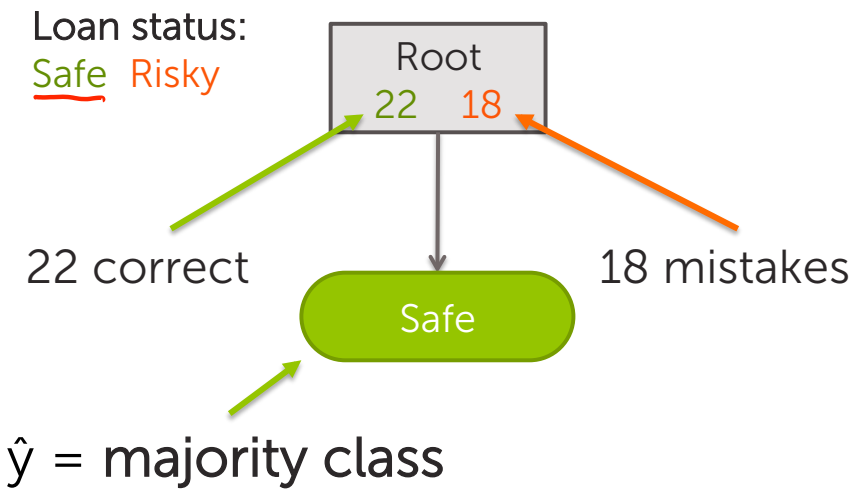

$$\text{Error} = \frac{18}{22+18} = 0.45$$

Tree	Classification error
(root)	0.45

Choice 1: Split on **credit history**?

Choice 1: Split on Credit

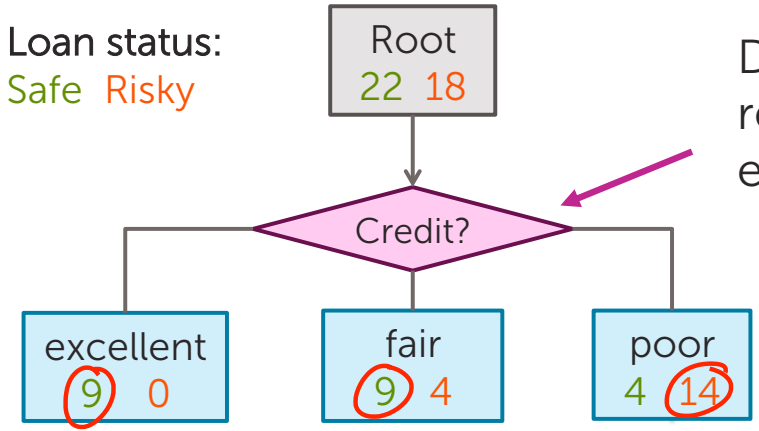

Does a **split on Credit** reduce classification error below 0.45?

How good is the split on **Credit**?

Choice 1: Split on **Credit**

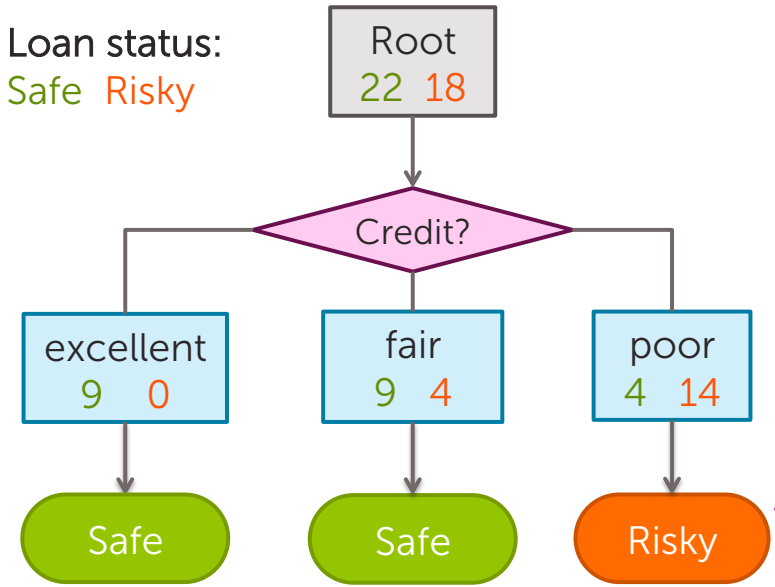

Step 1: For each intermediate node, set $\hat{y}$ = **majority value**

Split on **Credit**: Classification error

Choice 1: Split on Credit

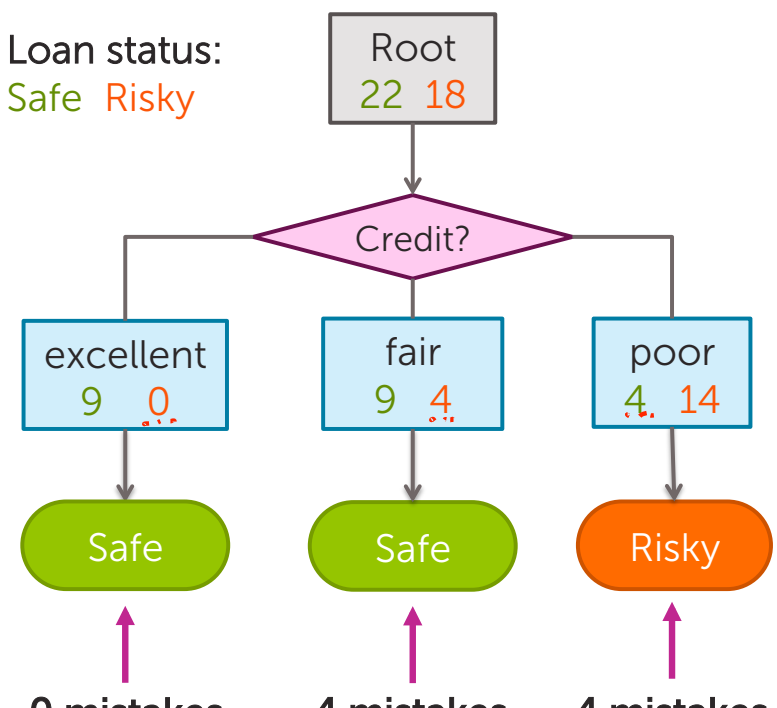

$$\text{Error} = \frac{4+4}{40} = 0.20$$

Tree	Classification error
(root)	0.45
Split on **credit**	0.2

Choice 2: Split on **Term**?

Choice 2: Split on Term

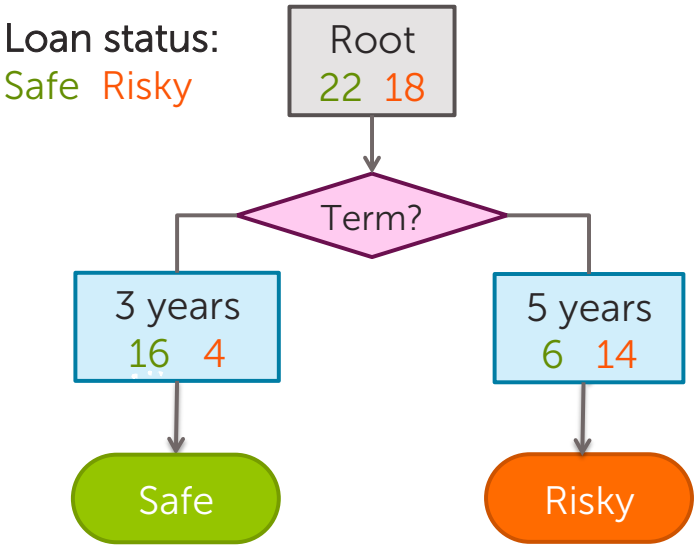

Evaluating the split on **Term**

Choice 2: Split on Term

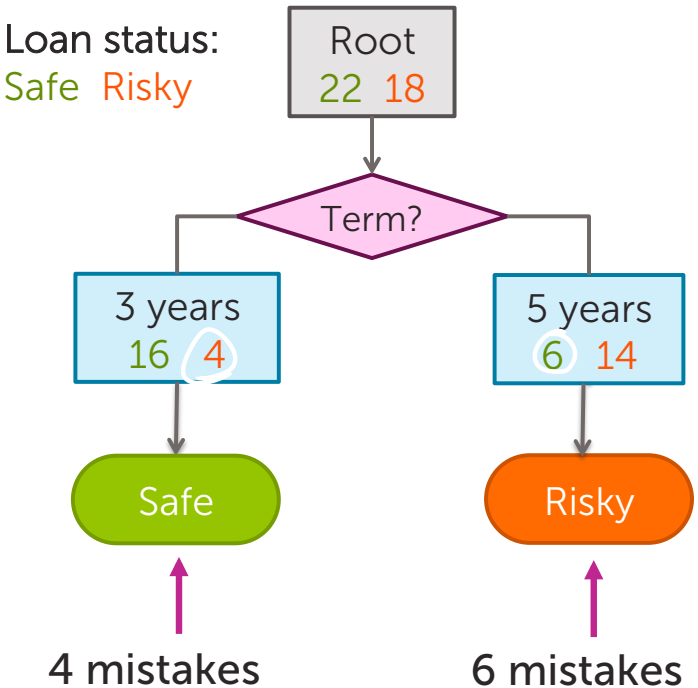

$$\text{Error} = \frac{4+6}{40} = 0.25$$

Tree	Classification error
(root)	0.45
Split on **credit**	0.2
Split on **term**	0.25

Choice 1 vs Choice 2

Tree	Classification error
(root)	0.45
split on **credit**	0.2
split on **loan term**	0.25

First split!

Choice 1: Split on Credit

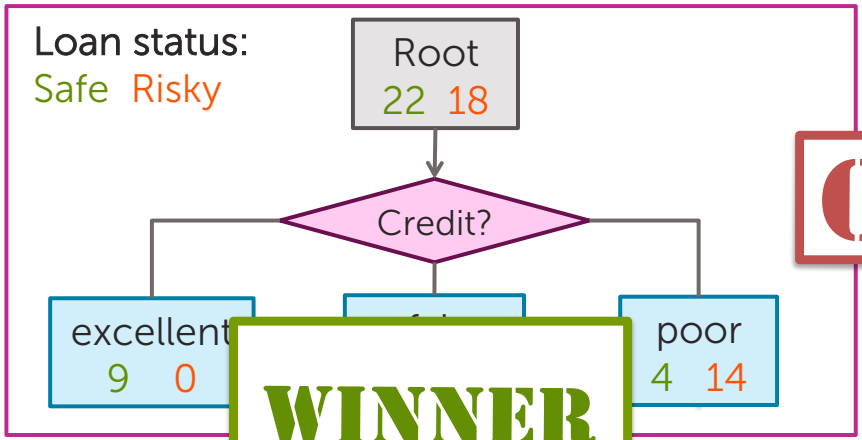

WINNER

OR

Choice 2: Split on Term

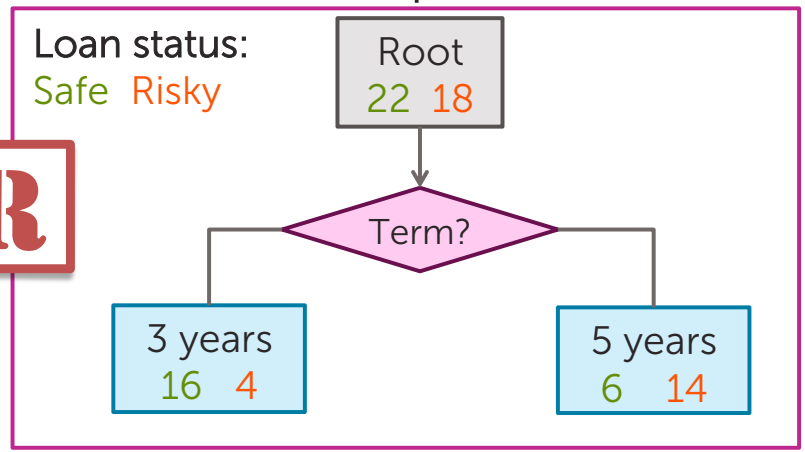

Feature split selection algorithm

- Given a subset of data M (a node in a tree)
- For each feature $h_i(x)$: credit, term, income
  1. Split data of M according to feature $h_i(x)$
  2. Compute classification error split
- Chose feature $h^*(x)$ with lowest classification error credit

Greedy decision tree learning

- Step 1: Start with an empty tree
- Step 2: Select a feature to split data
- For each split of the tree:
  - Step 3: If nothing more to, make predictions
  - Step 4: Otherwise, go to Step 2 & continue (recurse) on this split

Pick feature split leading to lowest classification error

Decision Tree Learning:
Recursion & Stopping conditions

Learn decision tree from data?

Credit	Term	Income	y
excellent	3 yrs	high	safe
fair	5 yrs	low	risky
fair	3 yrs	high	safe
poor	5 yrs	high	risky
excellent	3 yrs	low	risky
fair	5 yrs	low	safe
poor	3 yrs	high	risky
poor	5 yrs	low	safe
fair	3 yrs	high	safe

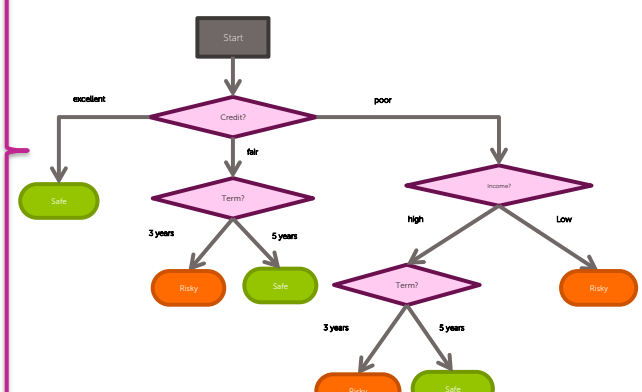

We've learned a decision stump, what next?

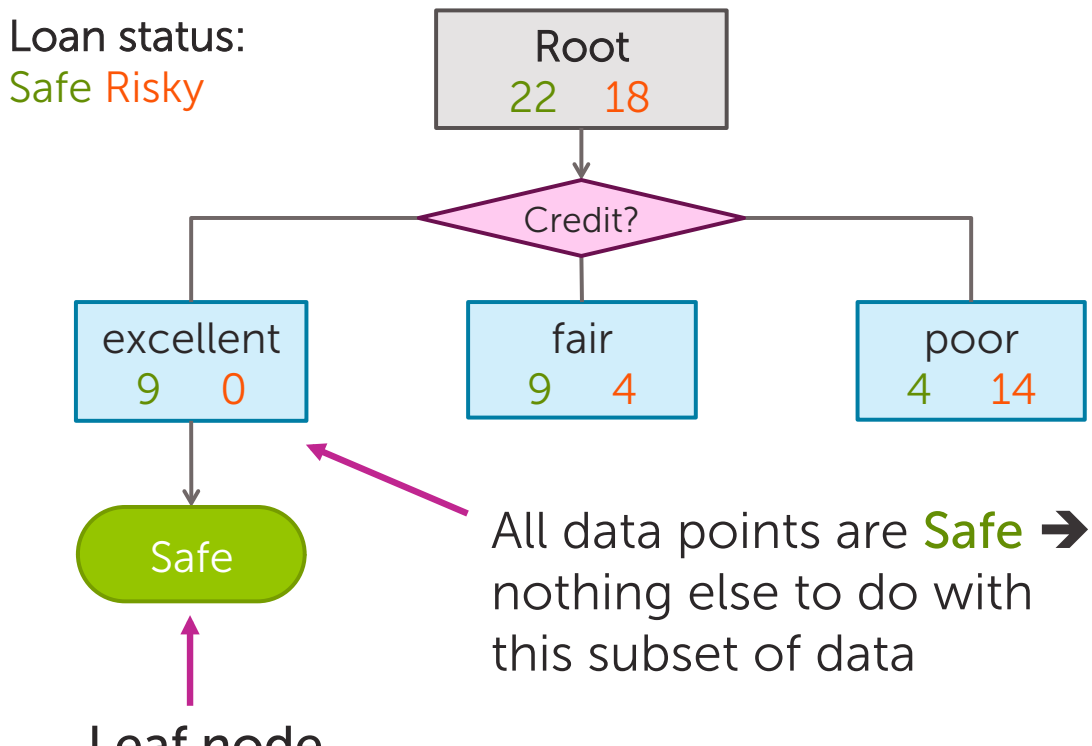

Tree learning = Recursive stump learning

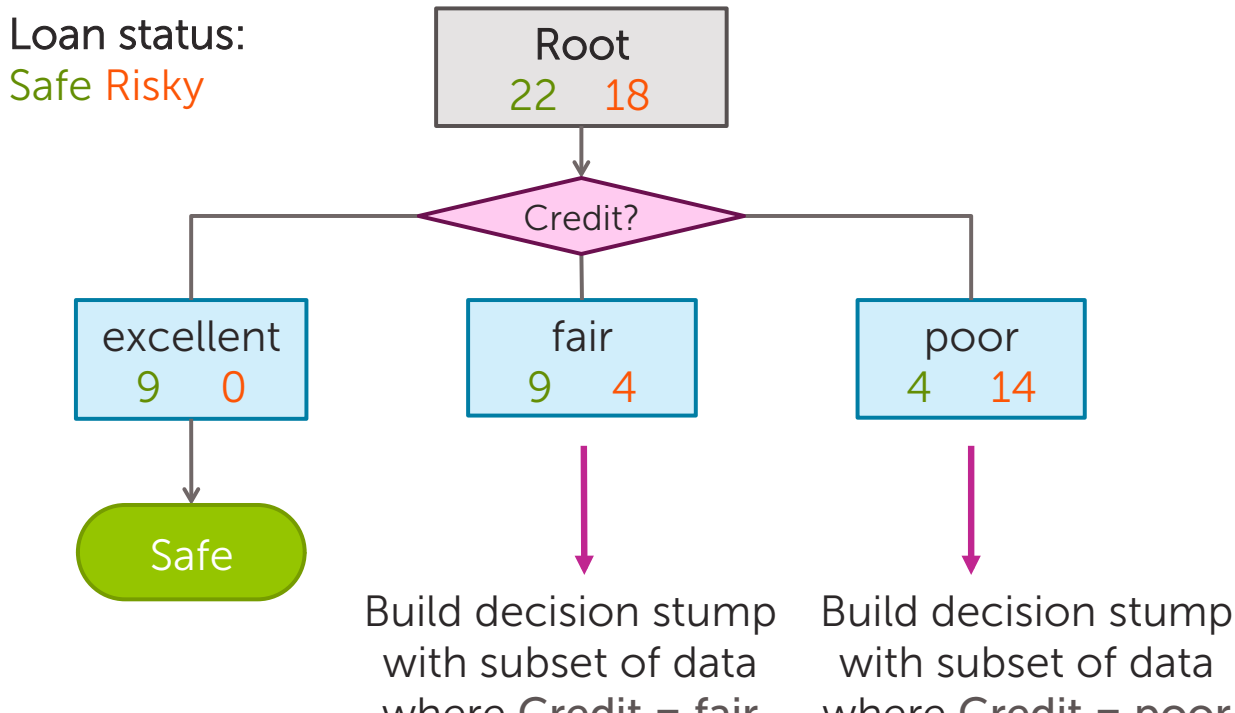

Second level

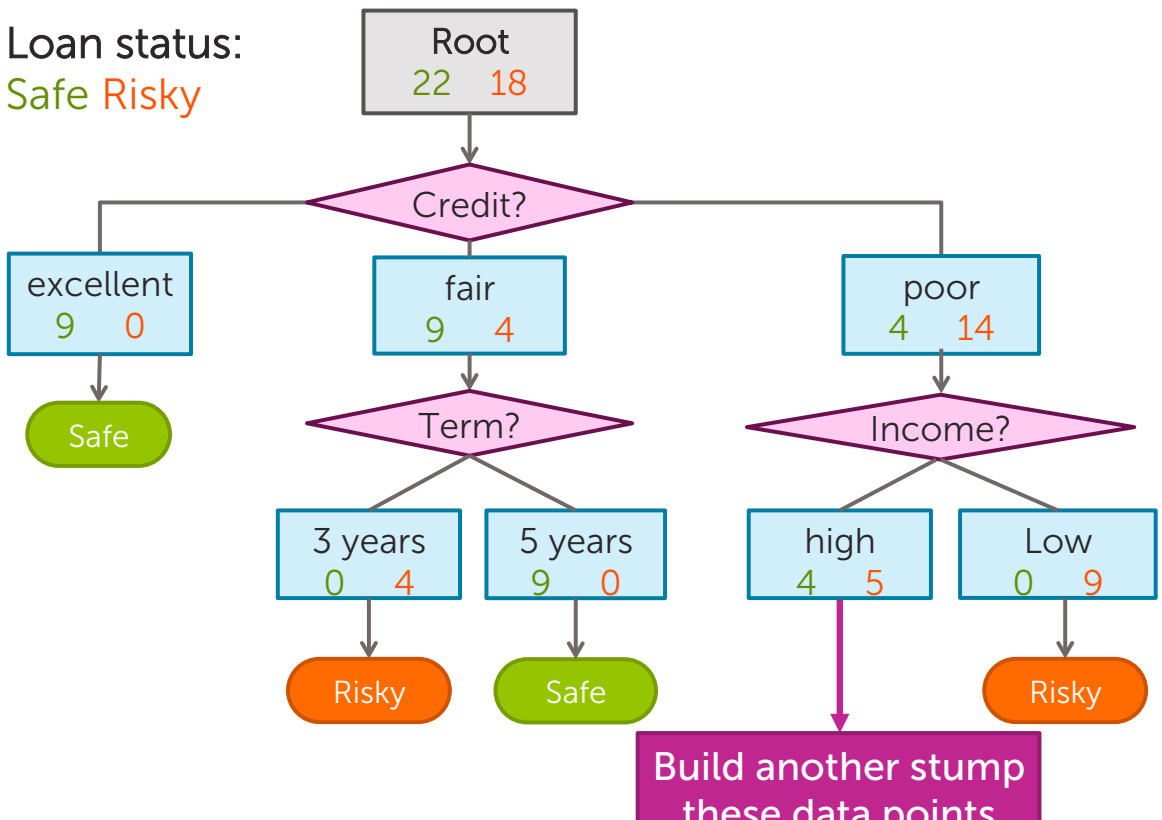

Final decision tree

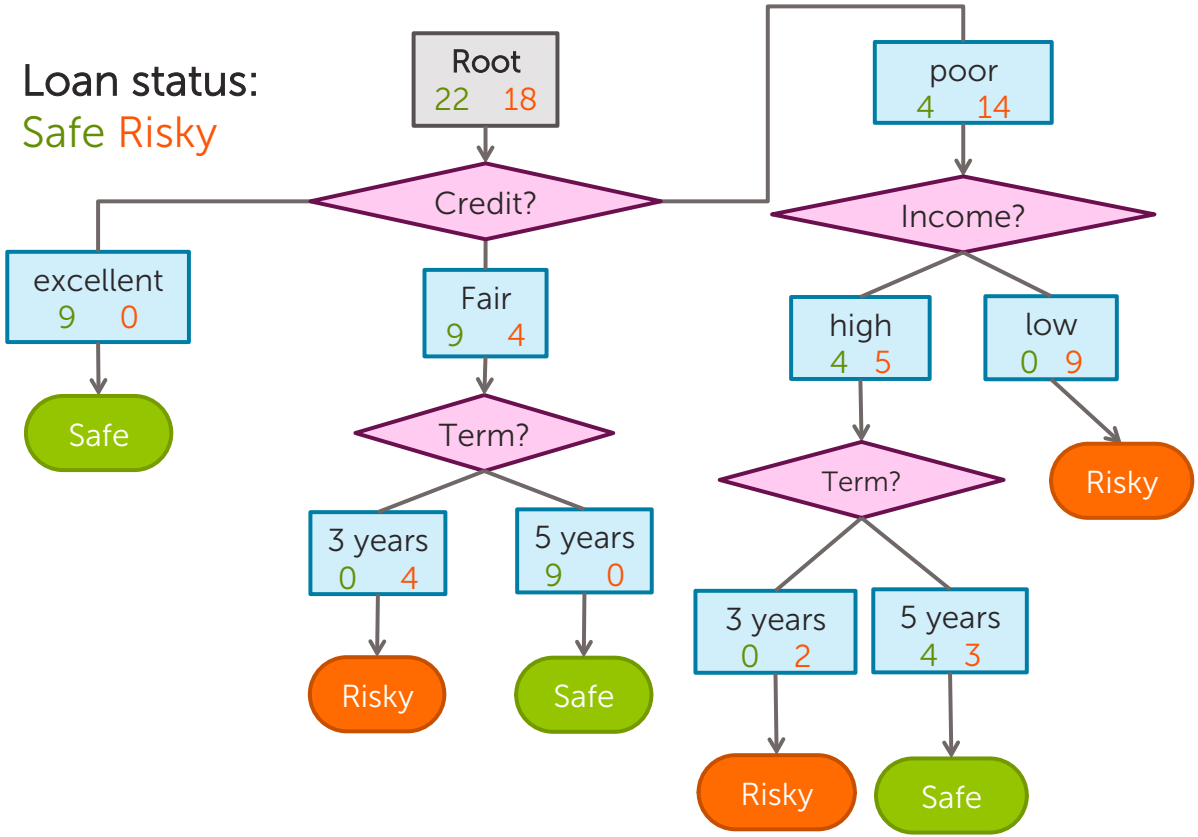

Simple greedy decision tree learning

Pick best feature to split on

Learn decision stump with this split

For each leaf of decision stump, recurse

When do we stop???

Stopping condition 1: All data agrees on y

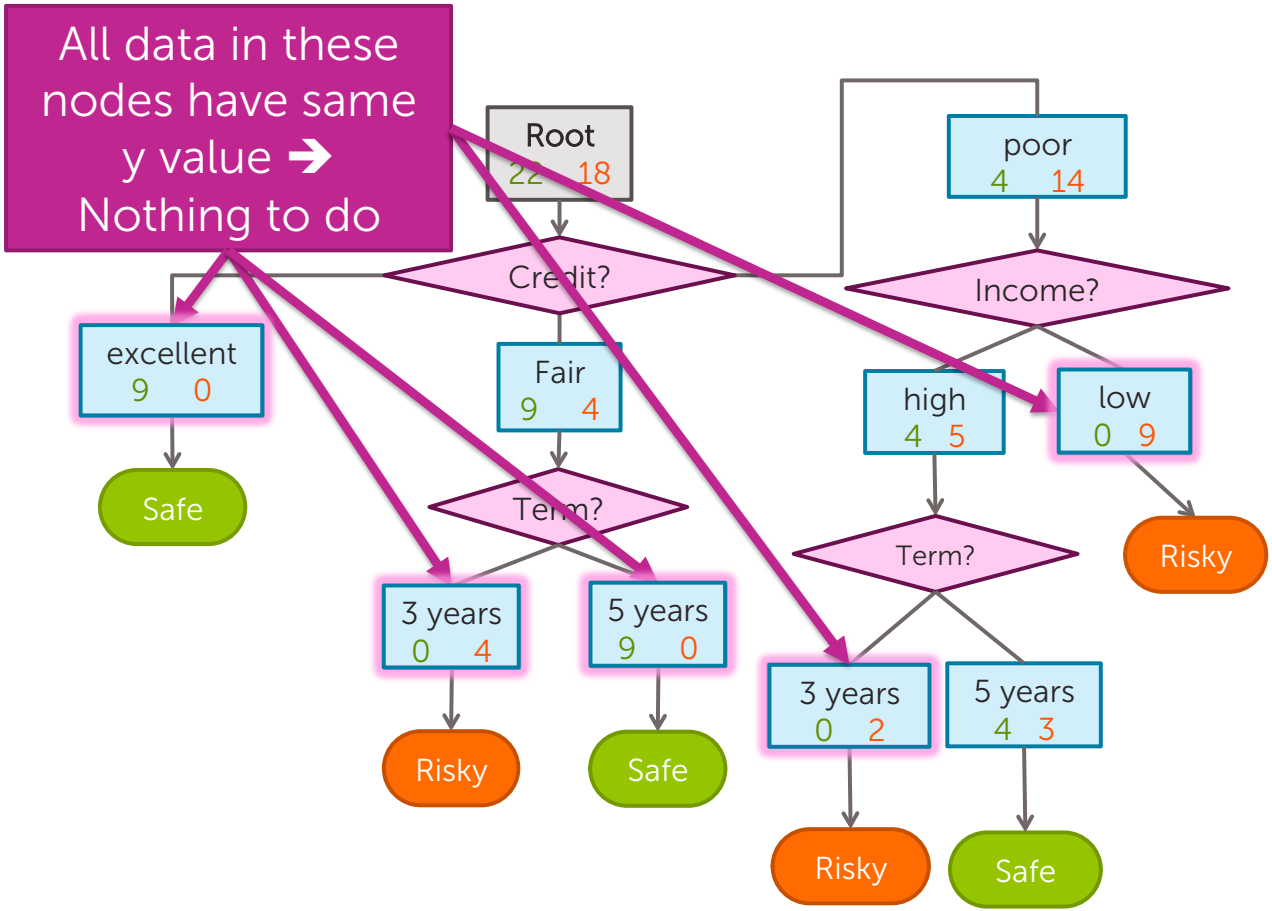

Stopping condition 2: Already split on all features

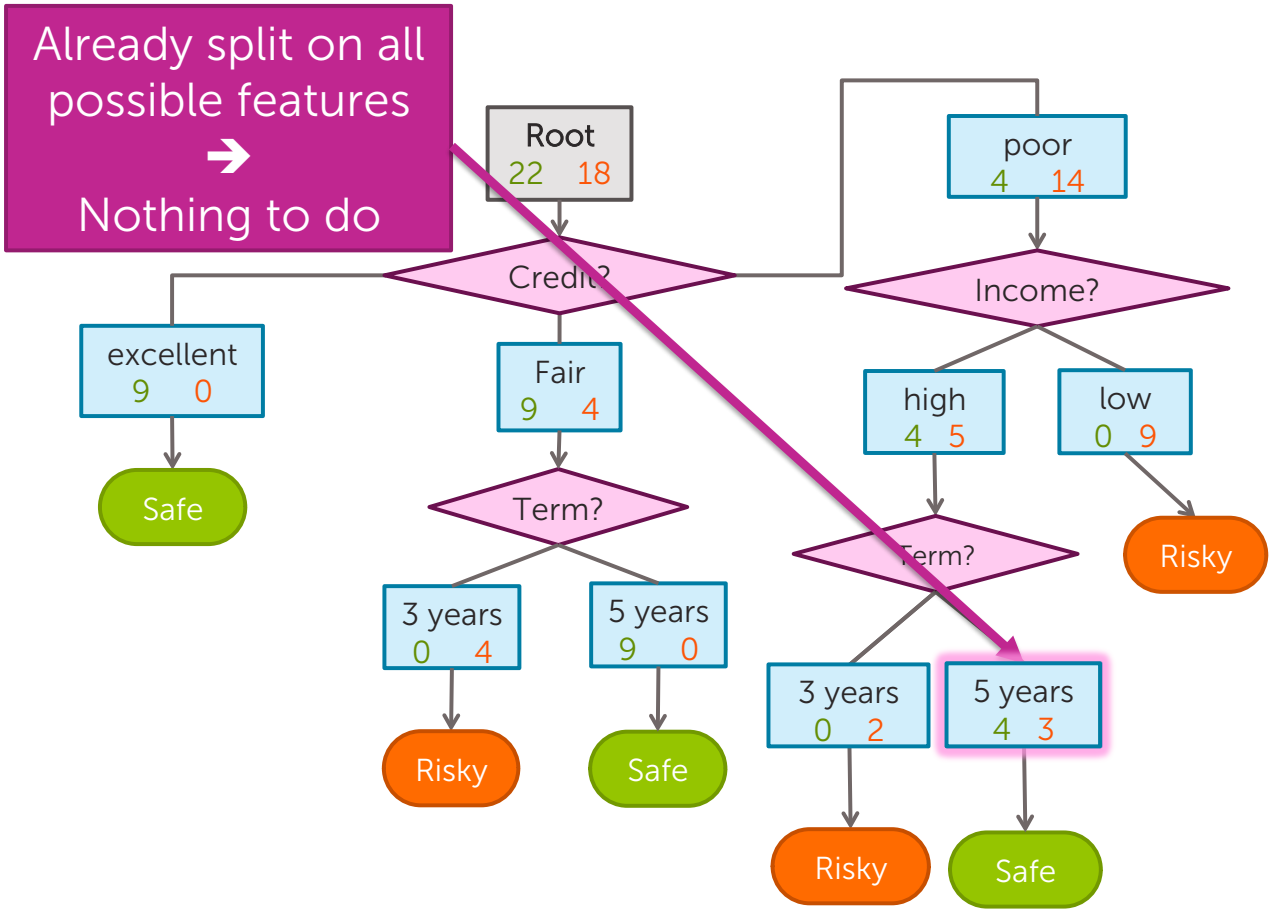

Greedy decision tree learning

- Step 1: Start with an empty tree
- Step 2: Select a feature to split data
- For each split of the tree:
  - Step 3: If nothing more to, make predictions
  - Step 4: Otherwise, go to Step 2 & continue (recurse) on this split

Pick feature split leading to lowest classification error

Stopping conditions 1 & 2

Recursion

Predictions with decision trees

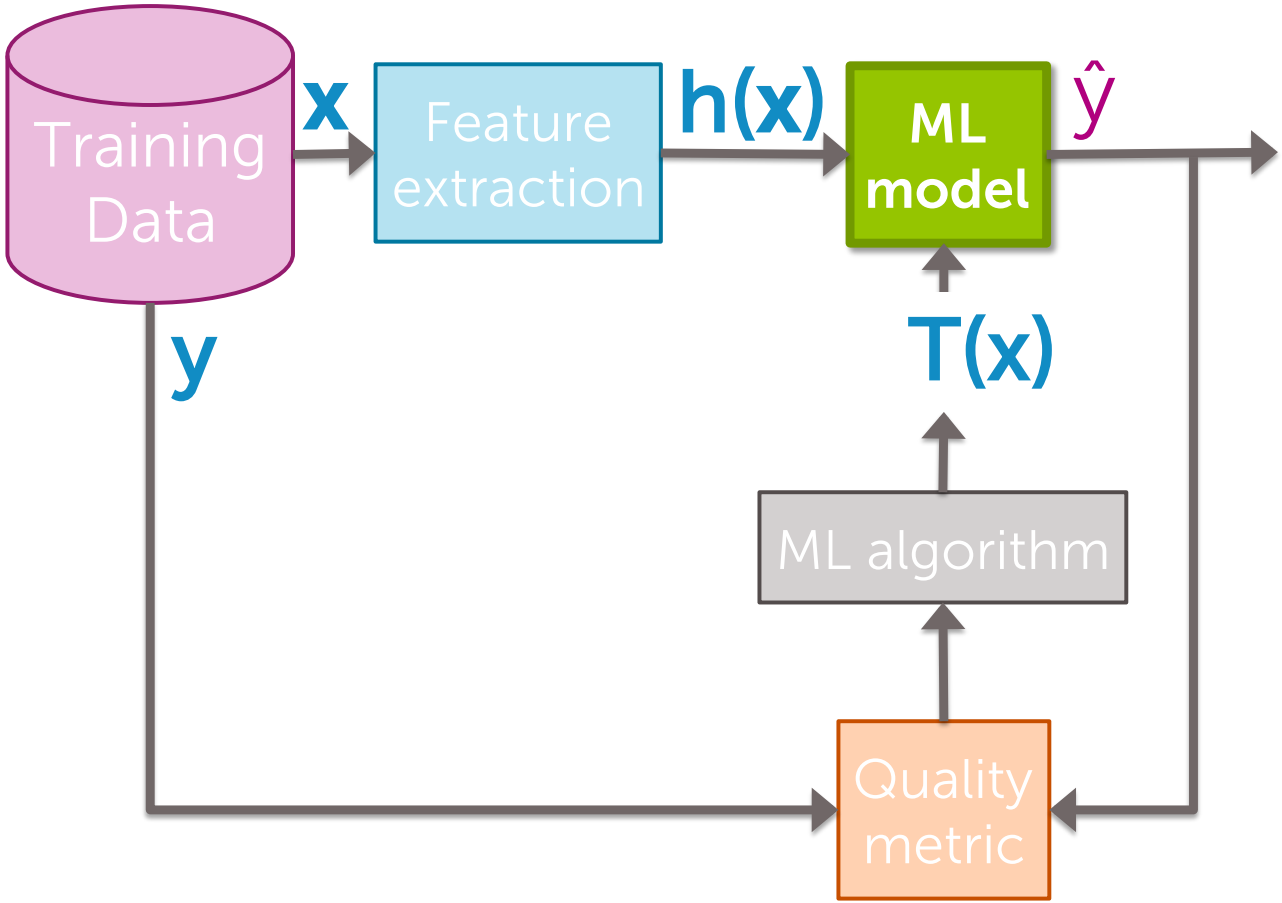
Training Data
x
Feature extraction
h(x)
ML model
ŷ
y
T(x)
ML algorithm
Quality metric

Decision tree model

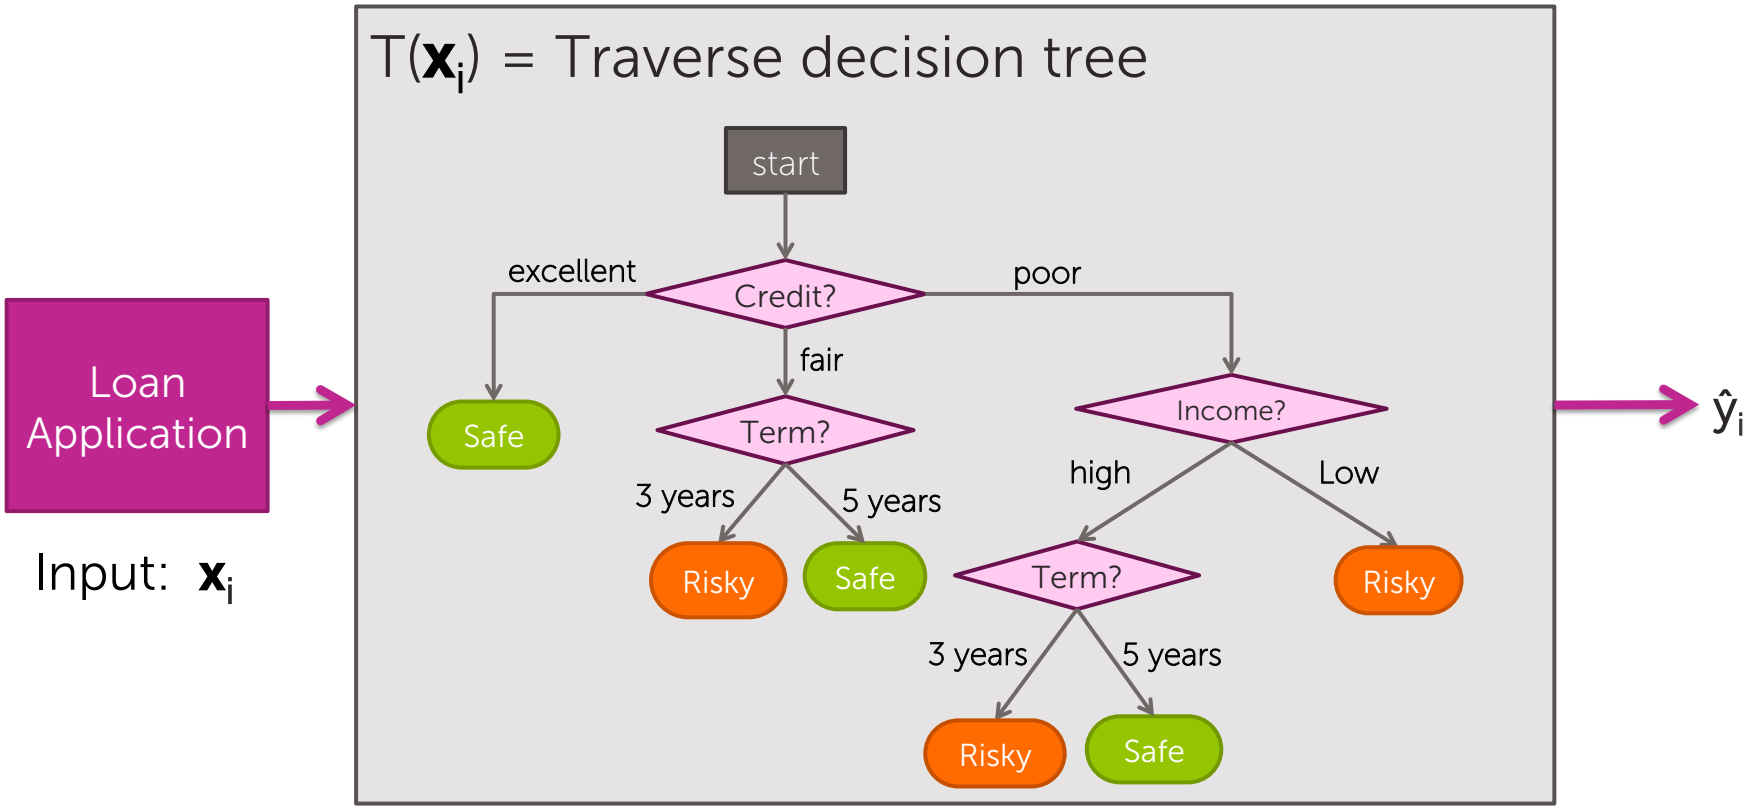

Traversing a decision tree

$\mathbf{x}_i$ = (Credit = poor, Income = high, Term = 5 years)

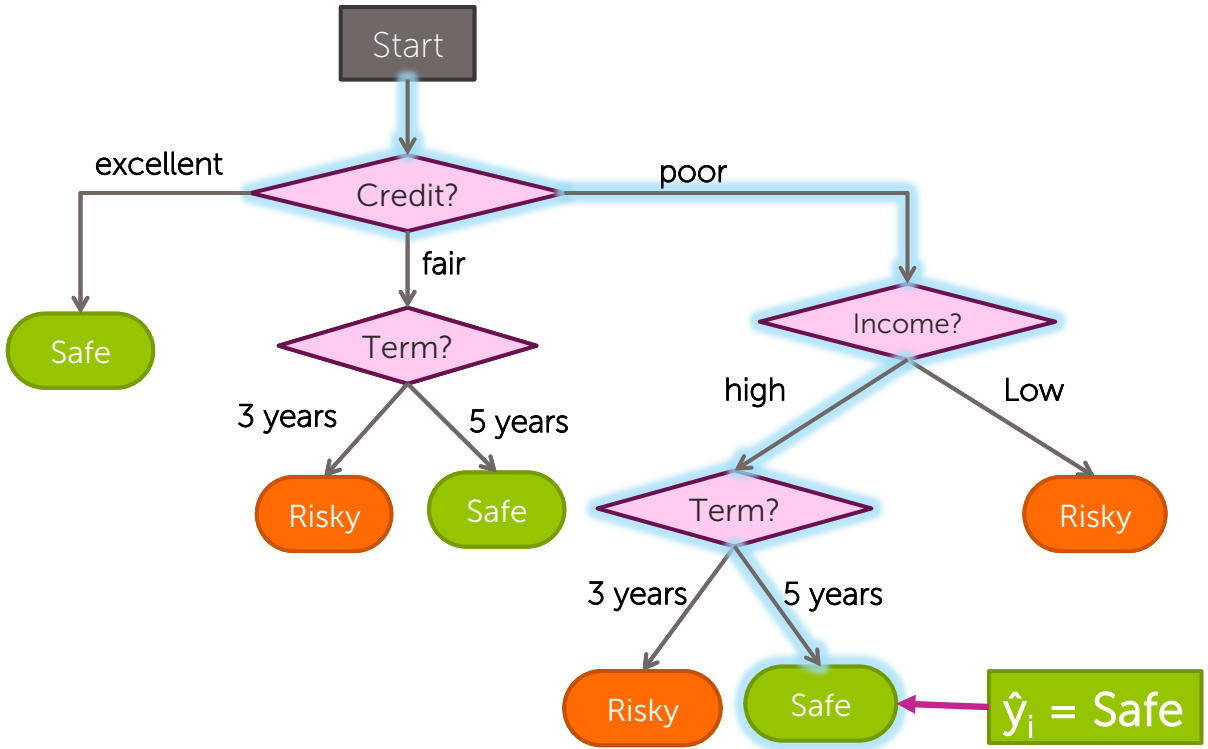

Decision tree prediction algorithm

predict(tree_node, input)

- If current tree_node is a leaf:
  - return majority class of data points in leaf
- else:
  - next_note = child node of tree_node whose feature value agrees with input
  - return predict(next_note, input)

Multiclass classification & predicting probabilities

Multiclass prediction

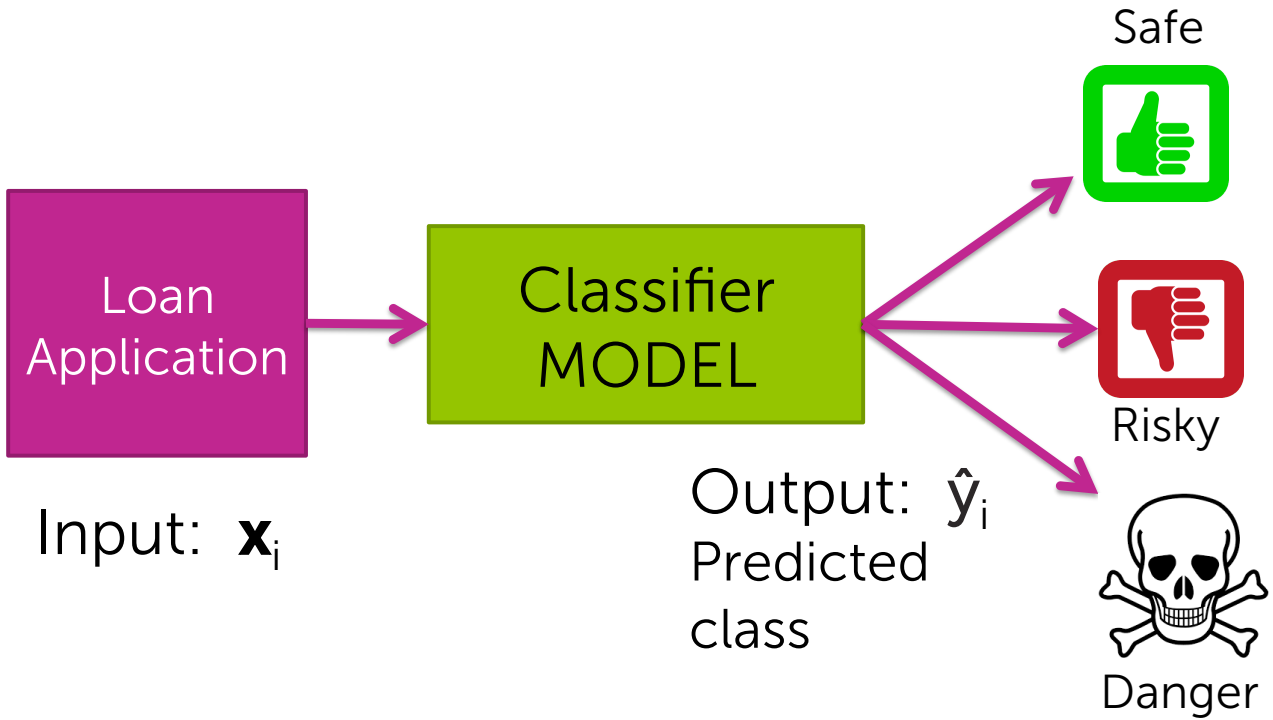

Input: $\mathbf{x}_i$

Output: $\hat{y}_i$
Predicted class

Multiclass decision stump

N = 40,
1 feature,
3 classes

Credit	y
excellent	safe
fair	risky
fair	safe
poor	danger
excellent	risky
fair	safe
poor	danger
poor	safe
fair	safe
...	...

Loan status:
Safe Risky Danger

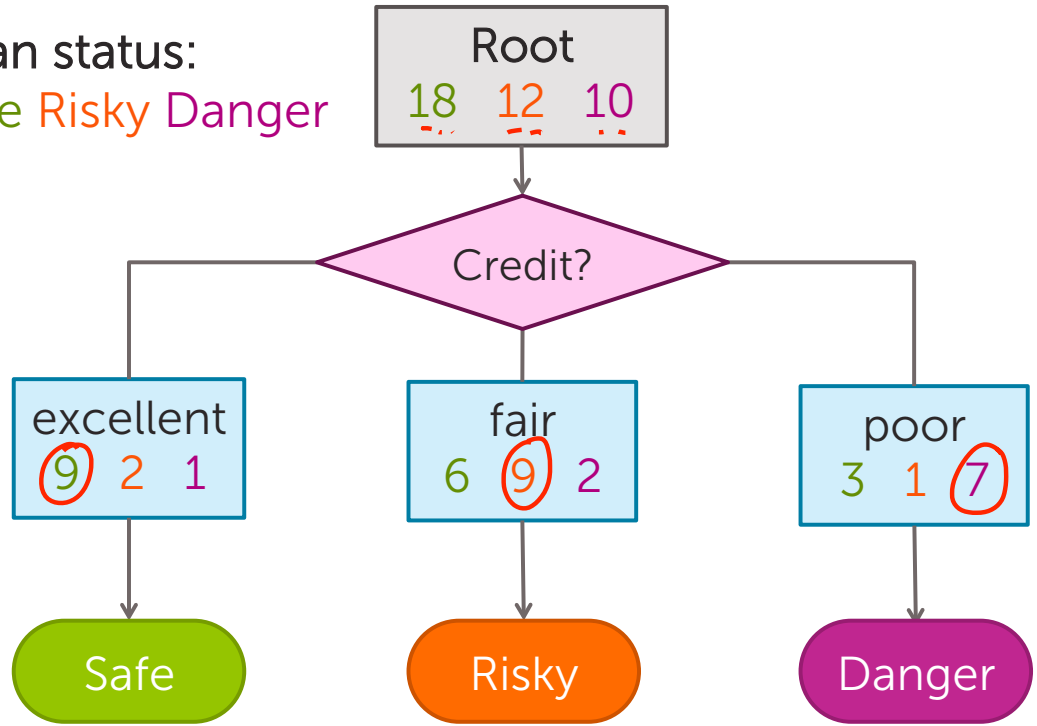

Predicting probabilities with decision trees

Loan status:
Safe Risky Danger

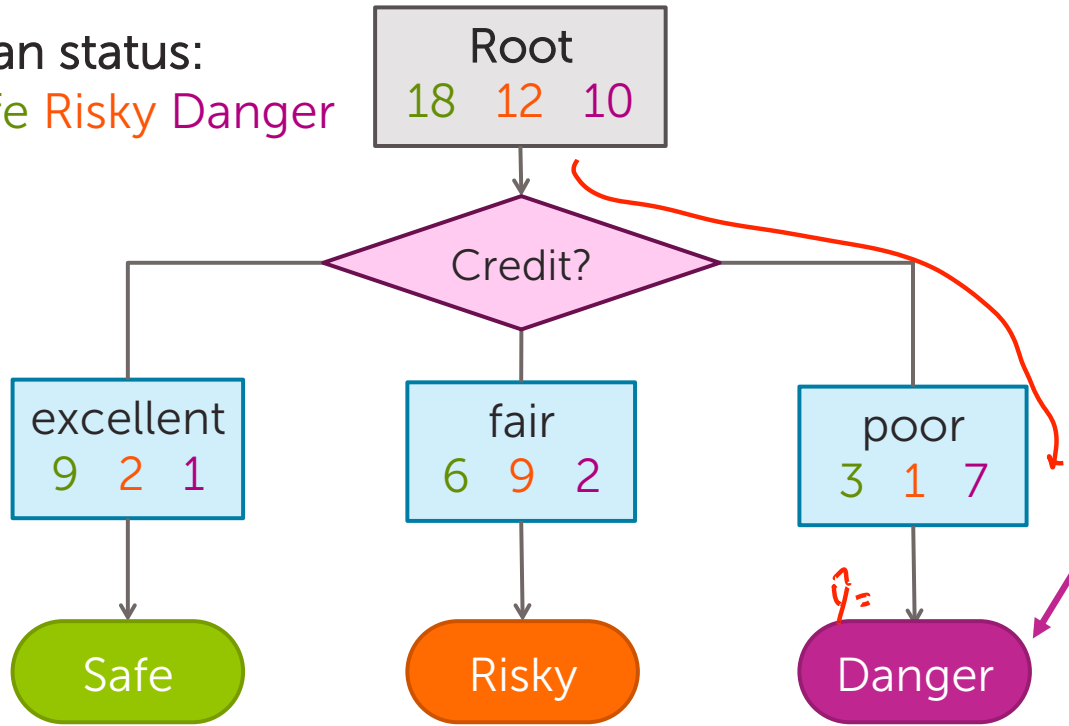

$$P(y = \text{danger} \mid \mathbf{x}) = \frac{7}{3 + 1 + 7} = 0.64$$

Decision tree learning:
Real valued features

How do we use real values inputs?

Income	Credit	Term	y
$105 K	excellent	3 yrs	Safe
$112 K	good	5 yrs	Risky
$73 K	fair	3 yrs	Safe
$69 K	excellent	5 yrs	Safe
$217 K	excellent	3 yrs	Risky
$120 K	good	5 yrs	Safe
$64 K	fair	3 yrs	Risky
$340 K	excellent	5 yrs	Safe
$60 K	good	3 yrs	Risky

Split on each numeric value?

Danger: May only contain one data point per node

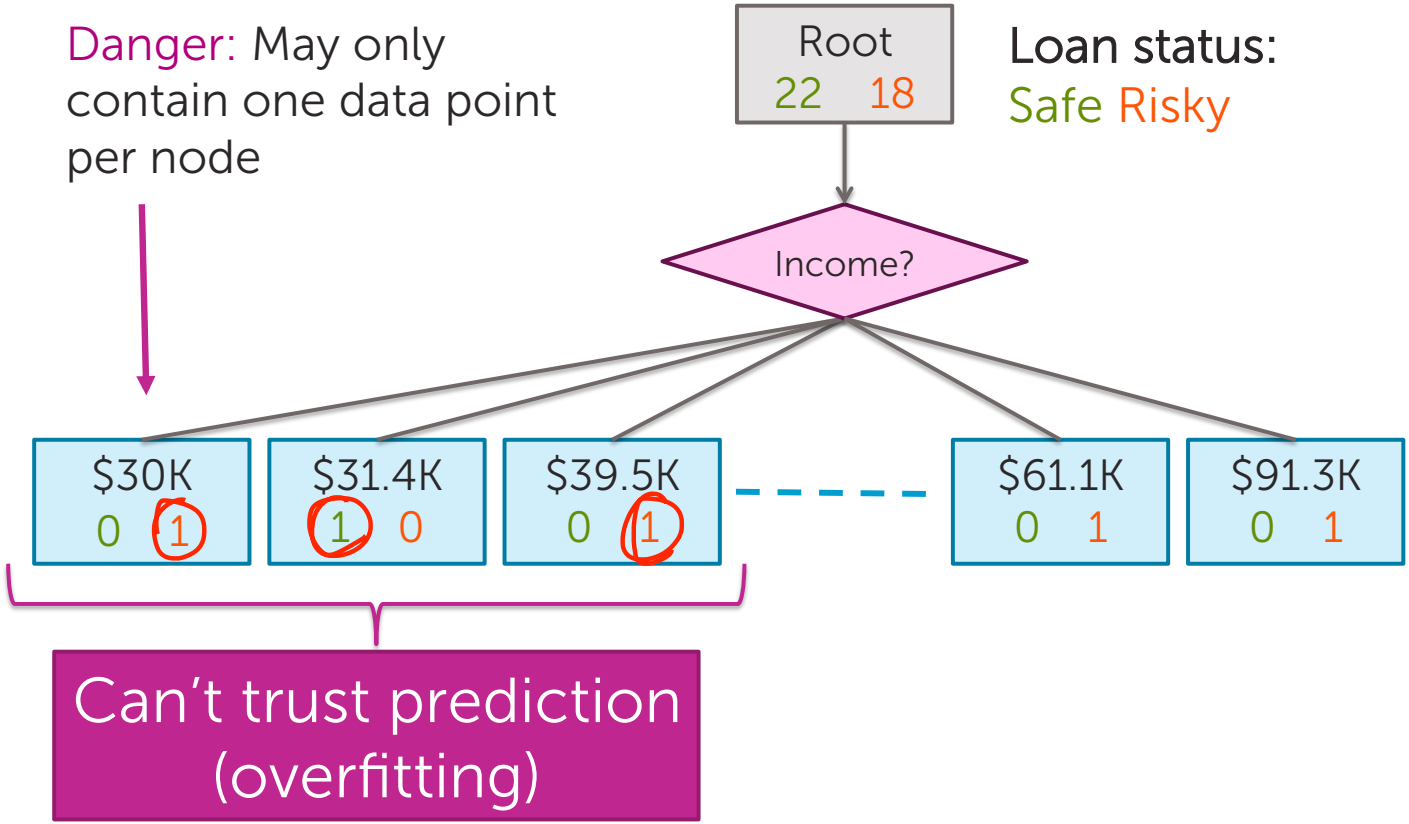

Alternative: Threshold split

Loan status:
Safe Risky

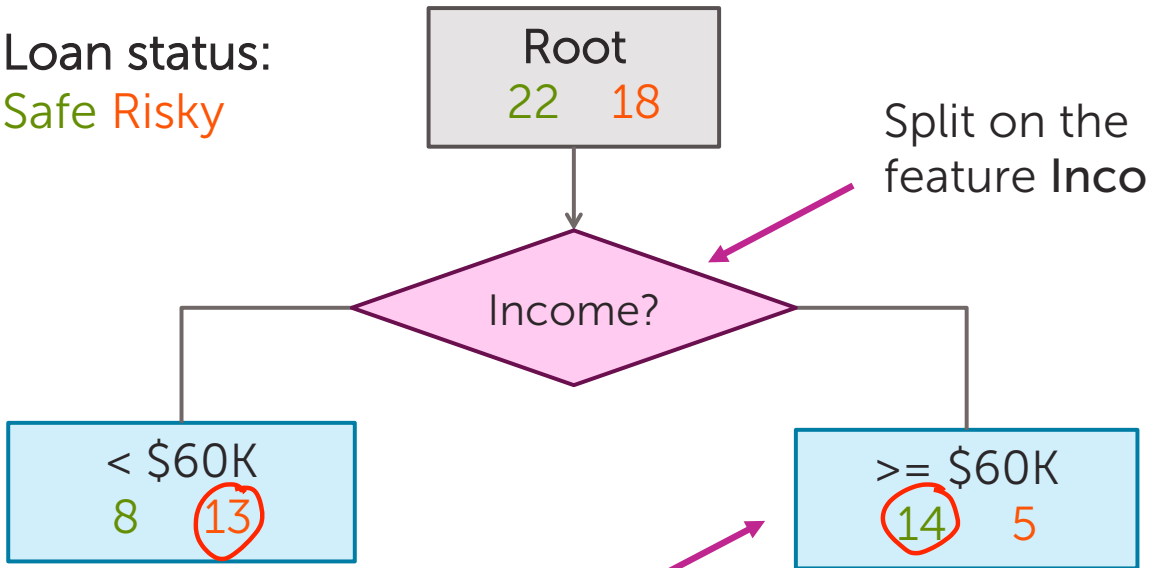

Split on the feature **Income**

Subset of data with Income >= $60K

Many data points ➔ lower chance of overfitting

Threshold splits in 1-D

Threshold split is the line
Income = $60K

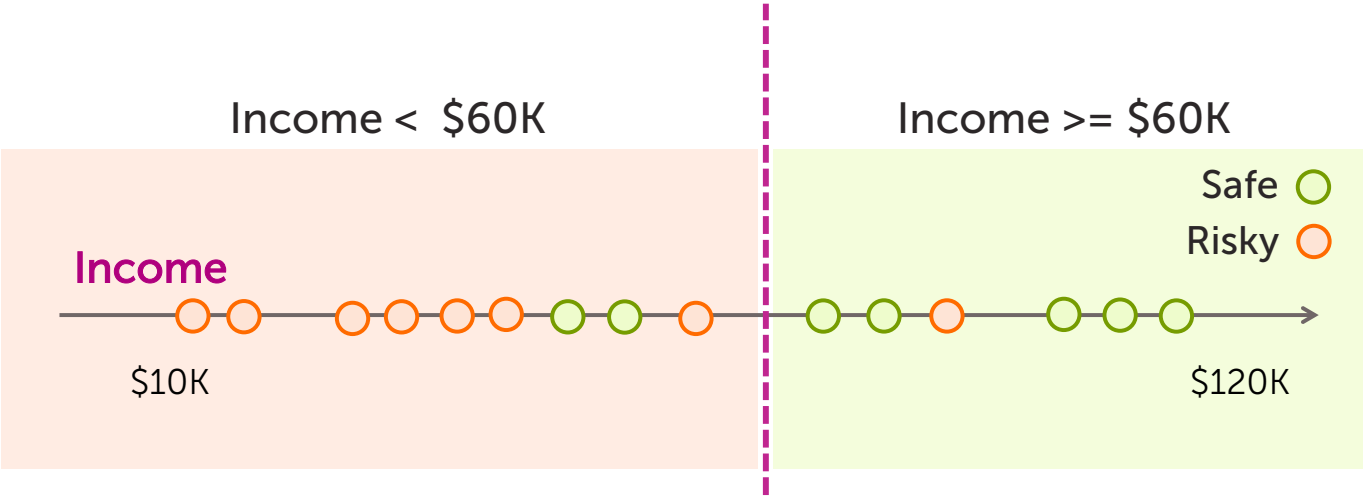

Visualizing the threshold split

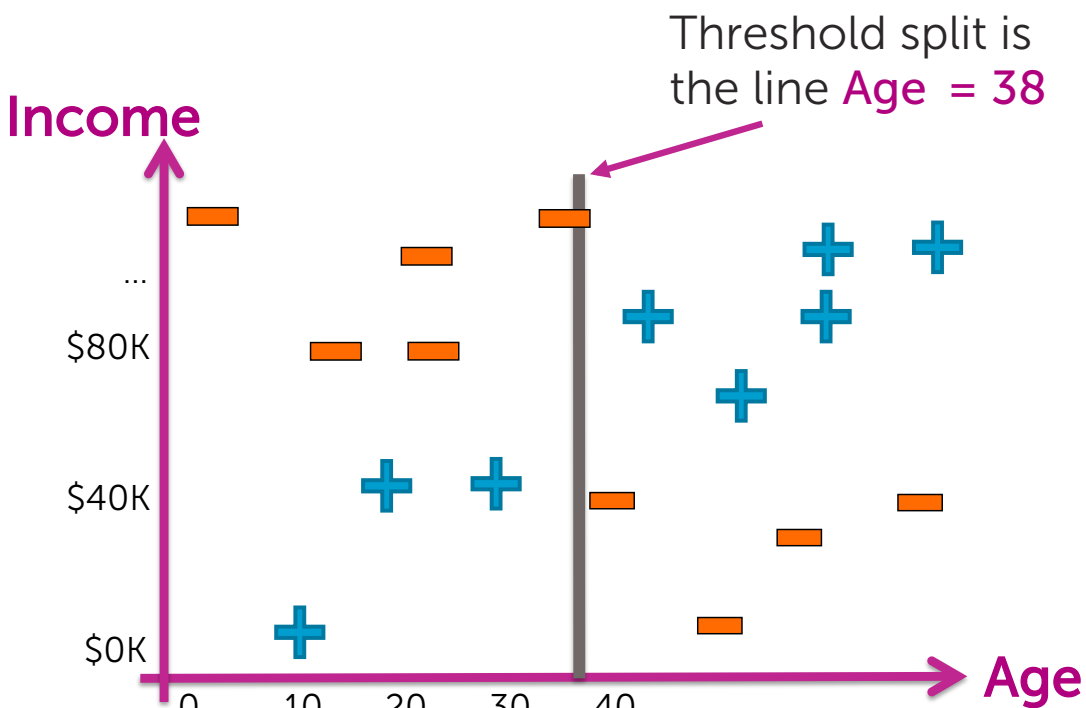

Split on Age >= 38

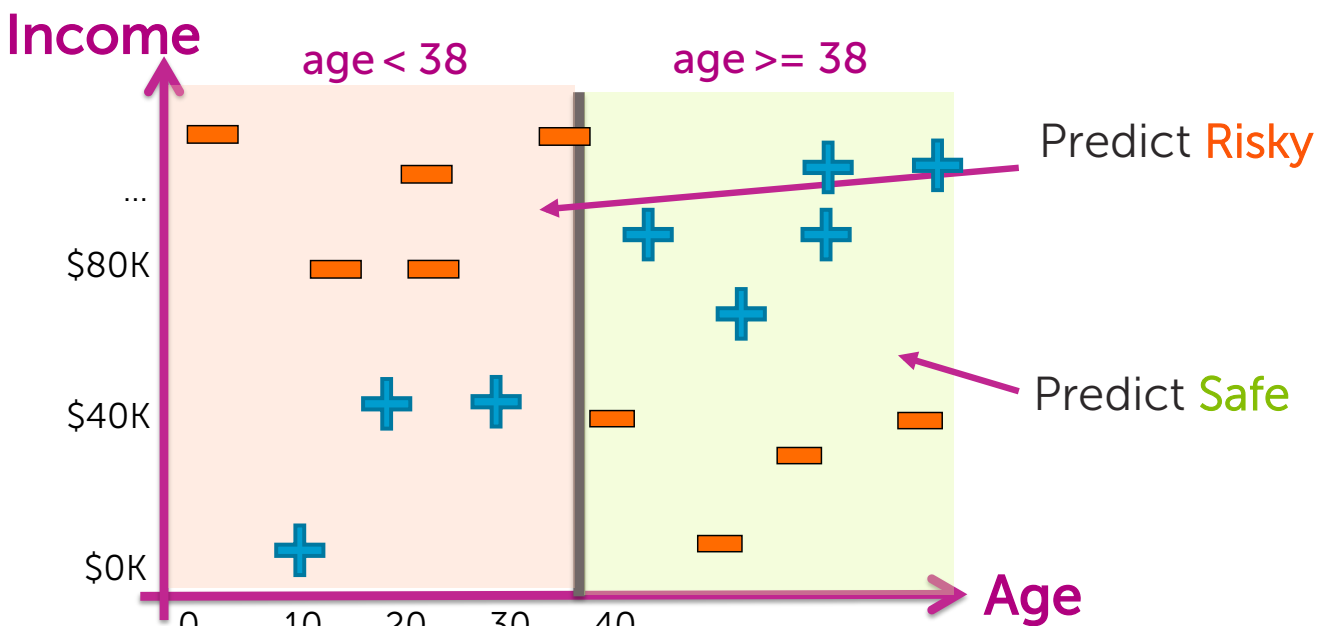

Depth 2: Split on Income >= $60K

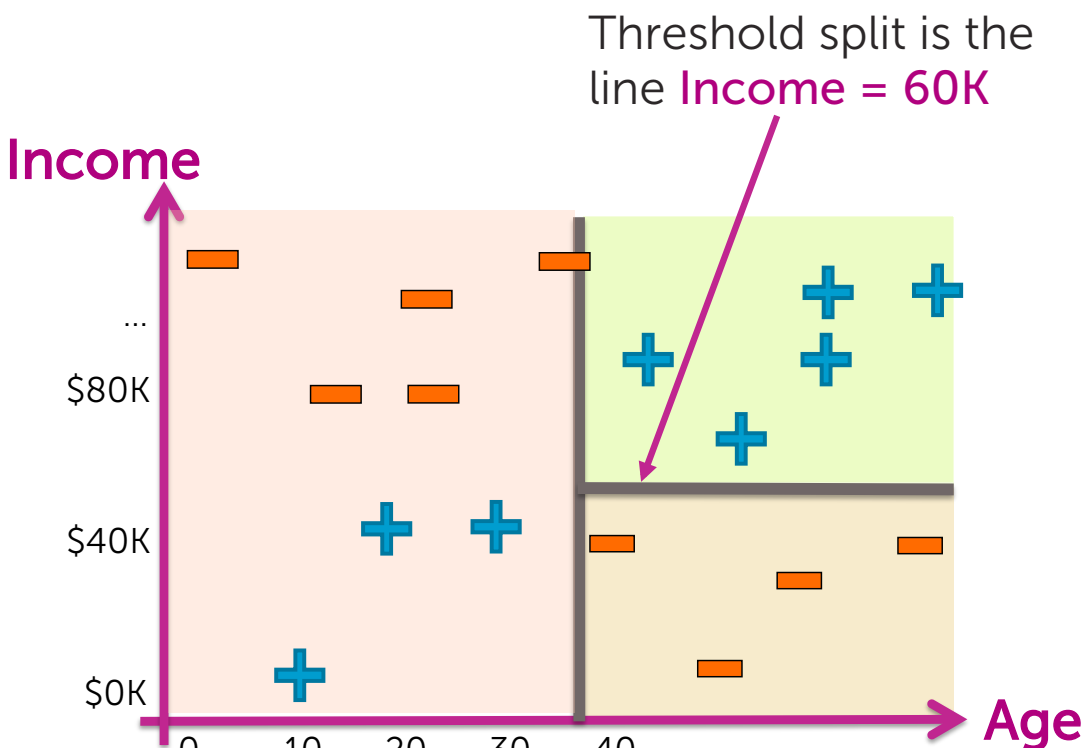

Each split partitions the 2-D space

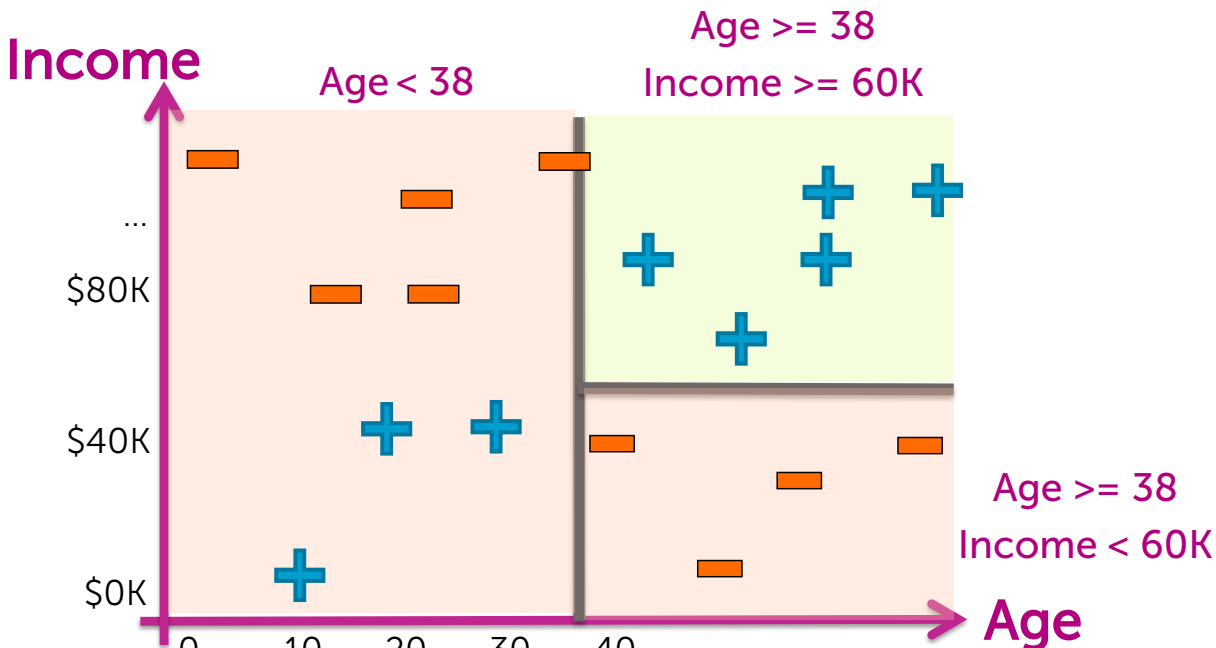

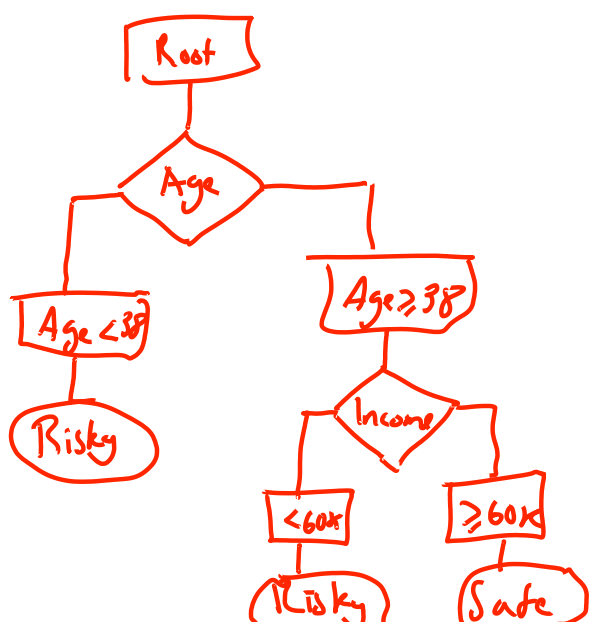

Finding the best threshold split

Infinite possible values of t

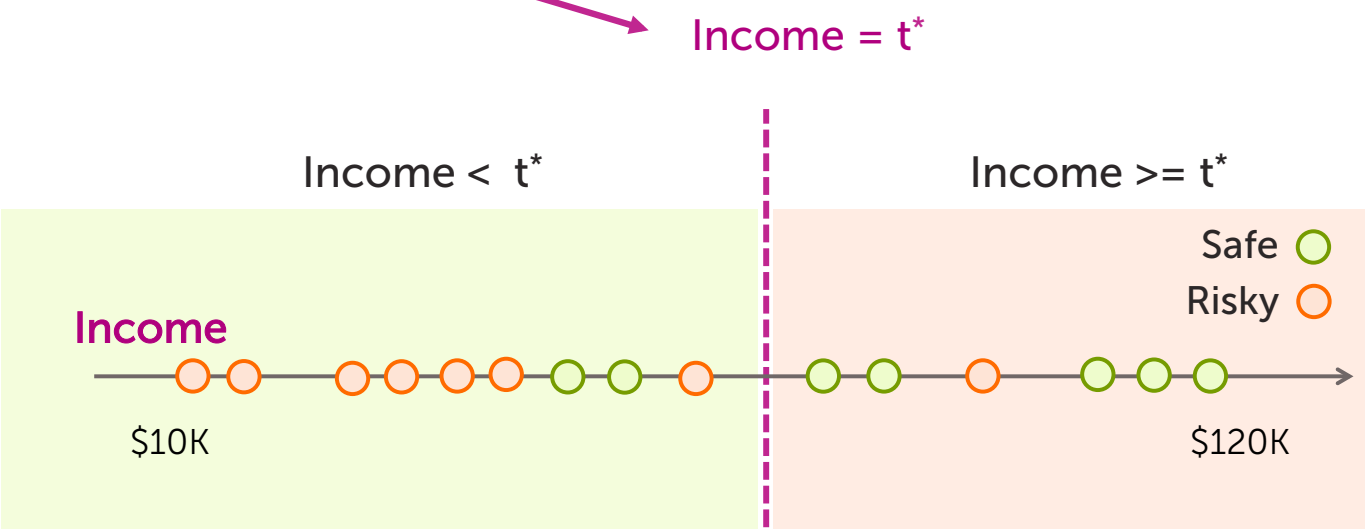

Consider a threshold between points

Same classification error for any threshold split between v_A and v_B

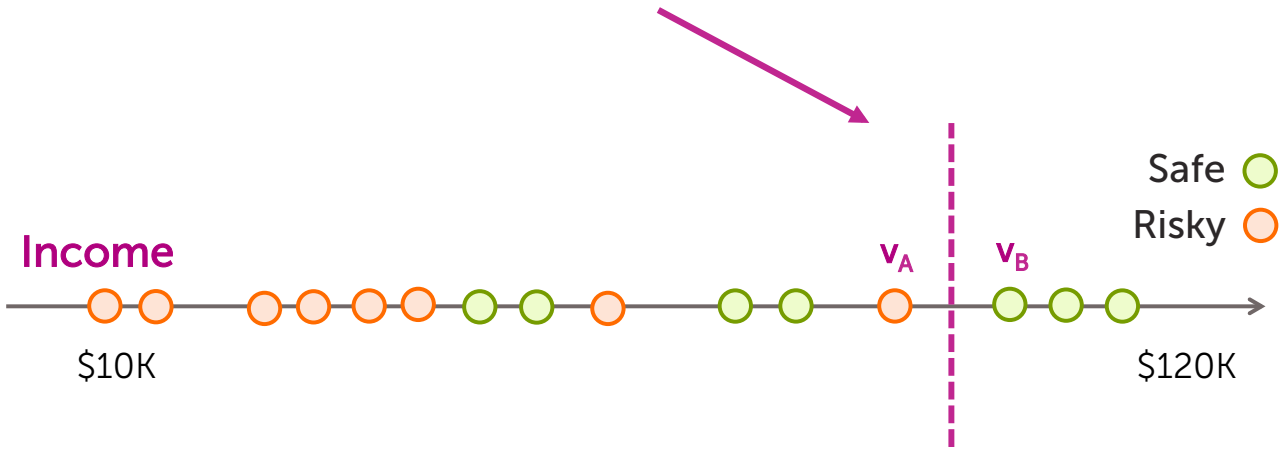

Only need to consider mid-points

Finite number of splits to consider

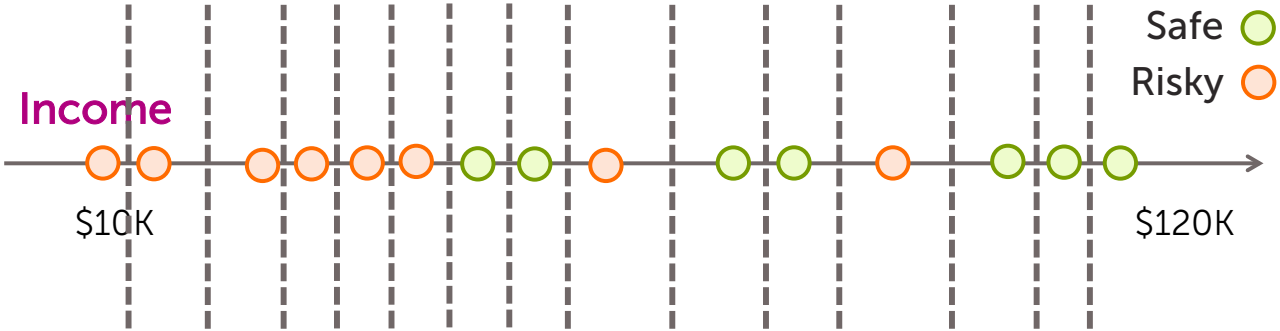

Threshold split selection algorithm

Income

- Step 1: Sort the values of a feature $h_j(\mathbf{x})$:

  Let $\{v_1, v_2, v_3, \ldots v_N\}$ denote sorted values

- Step 2:
  - For $i = 1 \ldots N-1$
    - Consider split $t_i = (v_i + v_{i+1}) / 2$
    - Compute classification error for treshold split $h_j(\mathbf{x}) >= t_i$
  - Chose the t^* with the lowest classification error

Decision trees vs logistic regression:
Example

Logistic regression

Feature	Value	Weight Learned
$h_0(\mathbf{x})$	1	0.22
$h_1(\mathbf{x})$	$\mathbf{x}[1]$	1.12
$h_2(\mathbf{x})$	$\mathbf{x}[2]$	-1.07

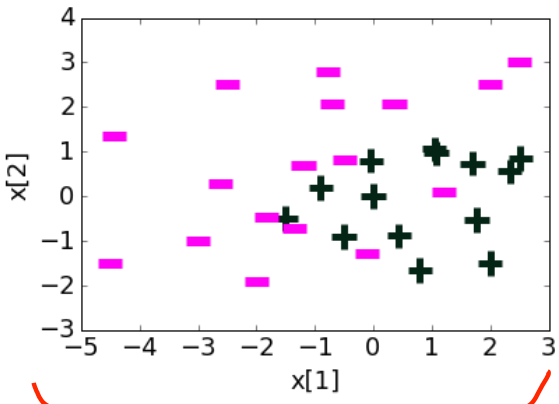

data

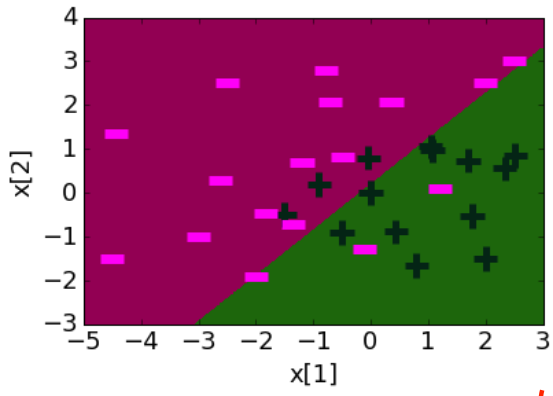

Depth 1: Split on **x**[1]

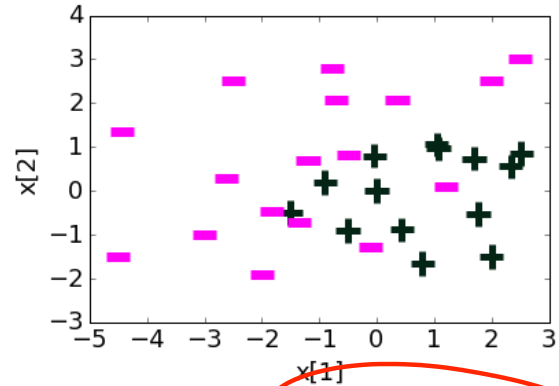

y values

- +

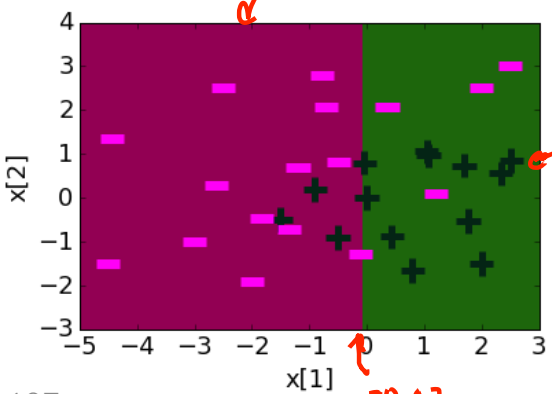

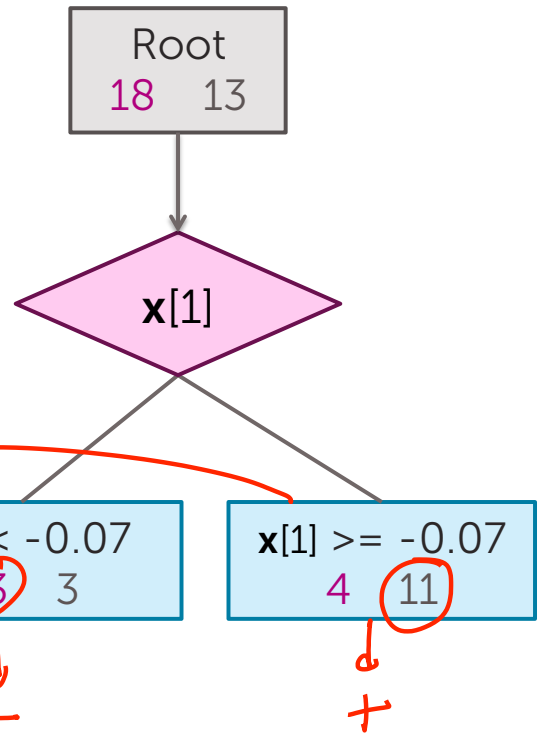

Depth 2

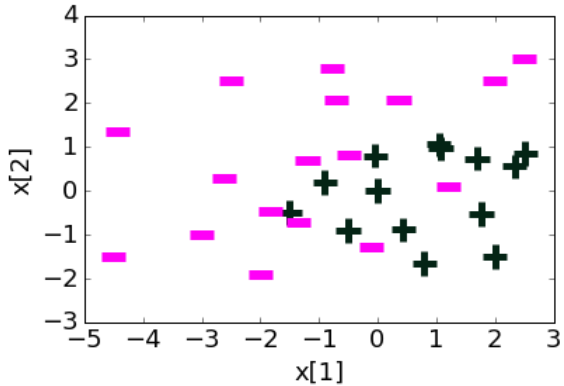
x[2]
x[1]

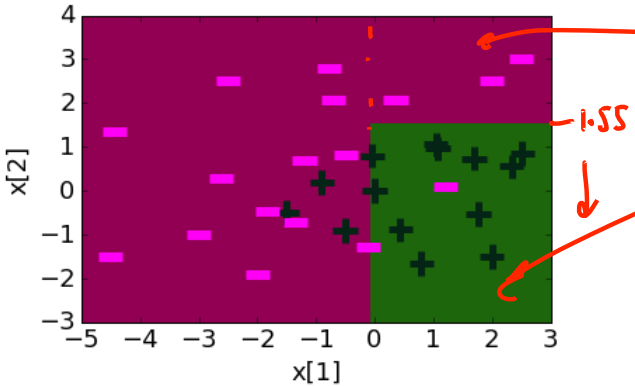
x[2]
x[1]
1.55

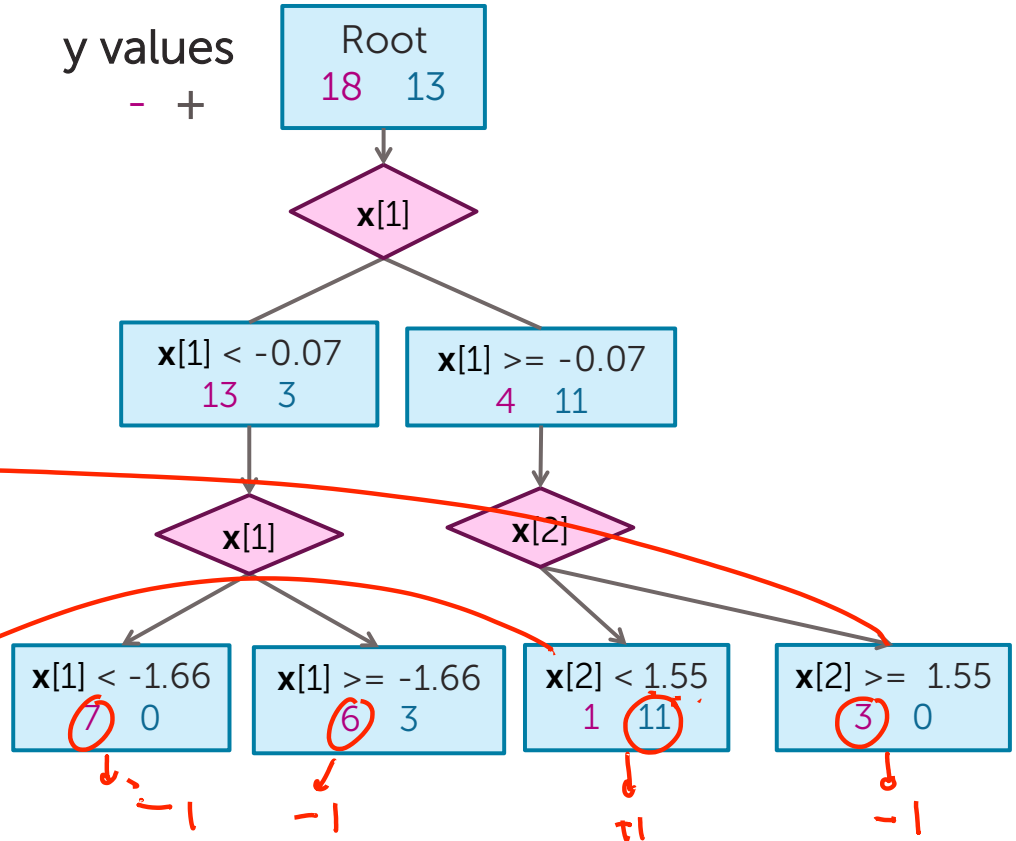
y values
- +
Root
18 13
x[1]
x[1] < -0.07
13 3
x[1] >= -0.07
4 11
x[1]
x[2]
x[1] < -1.66
7 0
x[1] >= -1.66
6 3
x[2] < 1.55
1 11
x[2] >= 1.55
3 0
-1
-1
+1
-1

Threshold split caveat

For threshold splits, same feature can be used multiple times

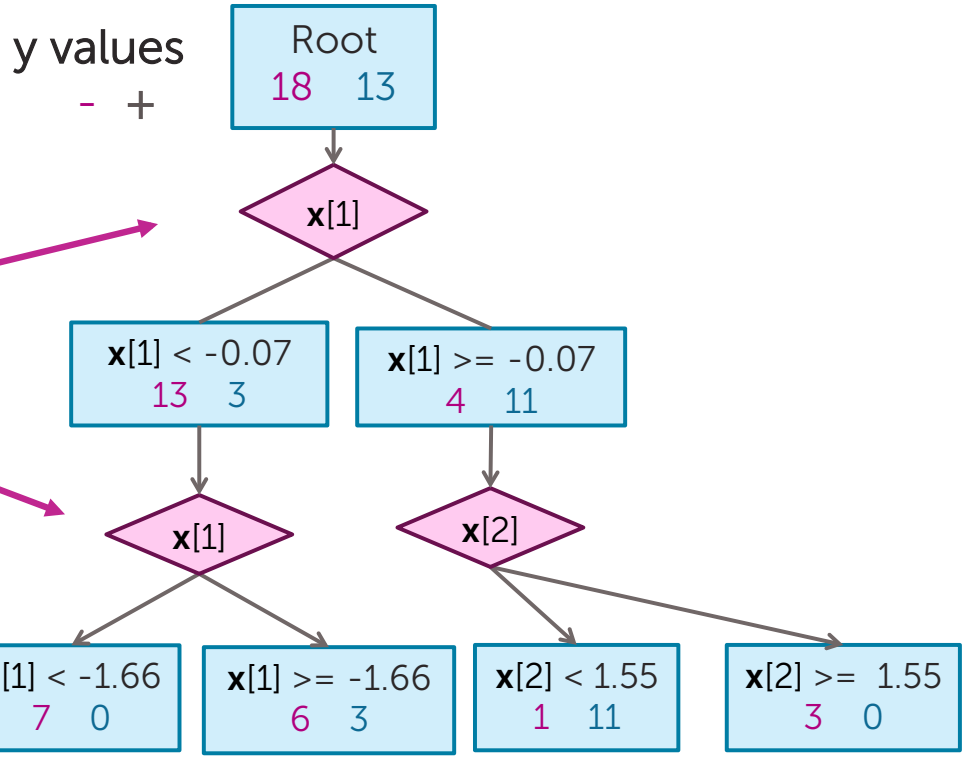

Decision boundaries

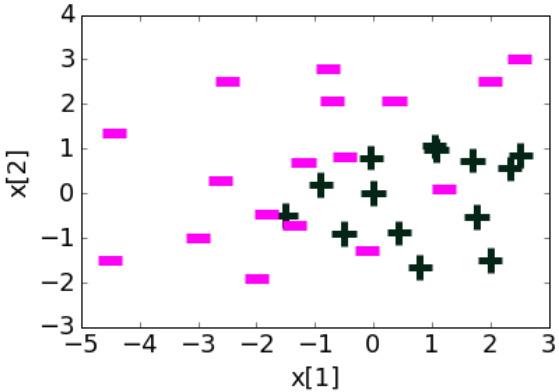

Depth 1

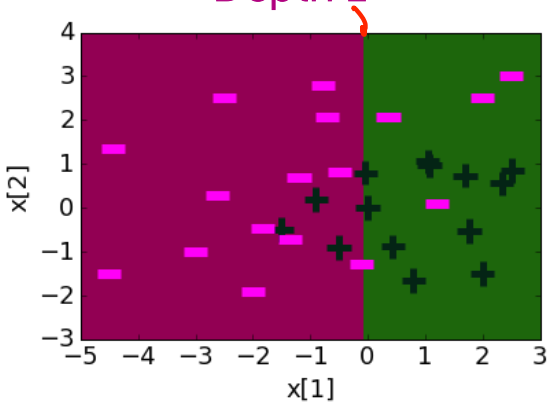

Depth 2

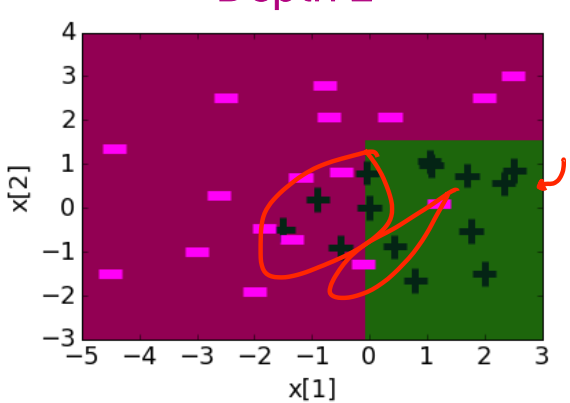

Depth 10

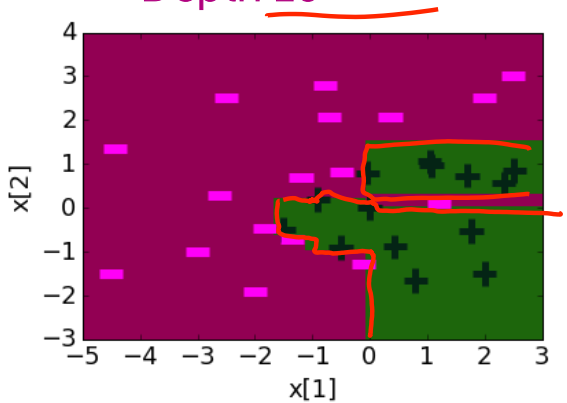

Comparing decision boundaries

Decision Tree

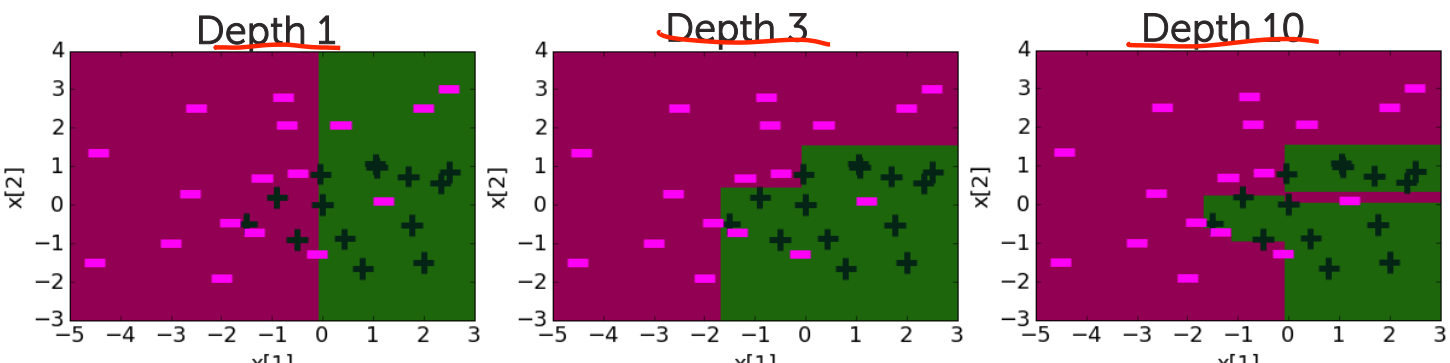

Logistic Regression

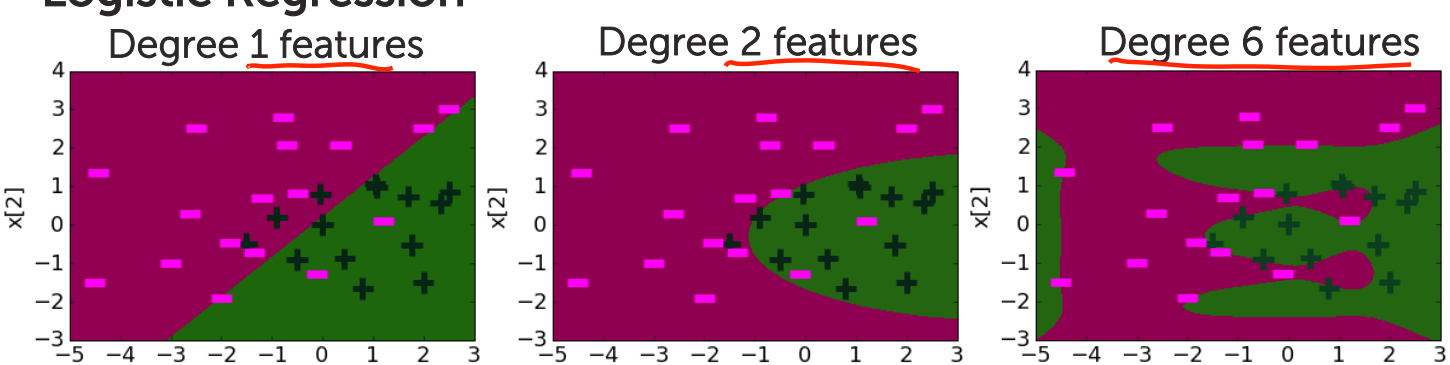

Summary of decision trees

What you can do now

- Define a decision tree classifier
- Interpret the output of a decision trees
- Learn a decision tree classifier using greedy algorithm
- Traverse a decision tree to make predictions
  - Majority class predictions
  - Probability predictions
  - Multiclass classification